

(automatic-prefix-caching)=

Automatic Prefix Caching

Introduction

Automatic Prefix Caching (APC in short) caches the KV cache of existing queries, so that a new query can directly reuse the KV cache if it shares the same prefix with one of the existing queries, allowing the new query to skip the computation of the shared part.

::{note}

Technical details on how vLLM implements APC can be found [\[here\]\(#design-automatic-prefix-caching\)](#).

:::

Enabling APC in vLLM

Set `enable_prefix_caching=True` in vLLM engine to enable APC. Here is an example:

```
```python
```

```
import time
```

```
from vllm import LLM, SamplingParams
```

```
A prompt containing a large markdown table. The table is randomly generated by GPT-4.
```

```
LONG_PROMPT = "You are a helpful assistant in recognizes the content of tables in markdown
format. Here is a table as follows.\n# Table\n" + ""
```

ID	Name	Age	Occupation	Country	Email	Phone Number	Address
1	John Doe	29	Engineer	USA	john.doe@example.com	555-1234	123 Elm St, Springfield, IL
2	Jane Smith	34	Doctor	Canada	jane.smith@example.com	555-5678	456 Oak St, Toronto, ON
3	Alice Johnson	27	Teacher	UK	alice.j@example.com	555-8765	789 Pine St, London, UK
4	Bob Brown	45	Artist	Australia	bob.b@example.com	555-4321	321 Maple St, Sydney, NSW
5	Carol White	31	Scientist	New Zealand	carol.w@example.com	555-6789	654 Birch St, Wellington, NZ
6	Dave Green	28	Lawyer	Ireland	dave.g@example.com	555-3456	987 Cedar St, Dublin, IE
7	Emma Black	40	Musician	USA	emma.b@example.com	555-1111	246 Ash St, New York, NY
8	Frank Blue	37	Chef	Canada	frank.b@example.com	555-2222	135 Spruce St, Vancouver, BC
9	Grace Yellow	50	Engineer	UK	grace.y@example.com	555-3333	864 Fir St, Manchester, UK
10	Henry Violet	32	Artist	Australia	henry.v@example.com	555-4444	753 Willow St, Melbourne, VIC
11	Irene Orange	26	Scientist	New Zealand	irene.o@example.com	555-5555	912 Poplar St, Auckland, NZ
12	Jack Indigo	38	Teacher	Ireland	jack.i@example.com	555-6666	159 Elm St, Cork, IE

13   Karen Red   41   Lawyer   USA   karen.r@example.com   555-7777   357
Cedar St, Boston, MA
14   Leo Brown   30   Chef   Canada   leo.b@example.com   555-8888   246
Oak St, Calgary, AB
15   Mia Green   33   Musician   UK   mia.g@example.com   555-9999   975
Pine St, Edinburgh, UK
16   Noah Yellow   29   Doctor   Australia   noah.y@example.com   555-0000   864
Birch St, Brisbane, QLD
17   Olivia Blue   35   Engineer   New Zealand   olivia.b@example.com   555-1212   753
Maple St, Hamilton, NZ
18   Peter Black   42   Artist   Ireland   peter.b@example.com   555-3434   912
Fir St, Limerick, IE
19   Quinn White   28   Scientist   USA   quinn.w@example.com   555-5656   159
Willow St, Seattle, WA
20   Rachel Red   31   Teacher   Canada   rachel.r@example.com   555-7878   357
Poplar St, Ottawa, ON
21   Steve Green   44   Lawyer   UK   steve.g@example.com   555-9090   753
Elm St, Birmingham, UK
22   Tina Blue   36   Musician   Australia   tina.b@example.com   555-1213   864
Cedar St, Perth, WA
23   Umar Black   39   Chef   New Zealand   umar.b@example.com   555-3435
975 Spruce St, Christchurch, NZ
24   Victor Yellow   43   Engineer   Ireland   victor.y@example.com   555-5657   246
Willow St, Galway, IE
25   Wendy Orange   27   Artist   USA   wendy.o@example.com   555-7879   135
Elm St, Denver, CO
26   Xavier Green   34   Scientist   Canada   xavier.g@example.com   555-9091   357

Oak St, Montreal, QC |

| 27 | Yara Red | 41 | Teacher | UK | yara.r@example.com | 555-1214 | 975

Pine St, Leeds, UK |

| 28 | Zack Blue | 30 | Lawyer | Australia | zack.b@example.com | 555-3436 | 135

Birch St, Adelaide, SA |

| 29 | Amy White | 33 | Musician | New Zealand | amy.w@example.com | 555-5658 |

159 Maple St, Wellington, NZ |

| 30 | Ben Black | 38 | Chef | Ireland | ben.b@example.com | 555-7870 | 246 Fir

St, Waterford, IE |

""

```
def get_generation_time(llm, sampling_params, prompts):
```

```
 # time the generation
```

```
 start_time = time.time()
```

```
 output = llm.generate(prompts, sampling_params=sampling_params)
```

```
 end_time = time.time()
```

```
 # print the output and generation time
```

```
 print(f"Output: {output[0].outputs[0].text}")
```

```
 print(f"Generation time: {end_time - start_time} seconds.")
```

```
set enable_prefix_caching=True to enable APC
```

```
llm = LLM(
```

```
 model='lmsys/longchat-13b-16k',
```

```
 enable_prefix_caching=True
```

```
)
```

```
sampling_params = SamplingParams(temperature=0, max_tokens=100)
```

```
Querying the age of John Doe
```

```
get_generation_time(
 llm,
 sampling_params,
 LONG_PROMPT + "Question: what is the age of John Doe? Your answer: The age of John Doe
is ",
)
```

```
Querying the age of Zack Blue
```

```
This query will be faster since vllm avoids computing the KV cache of LONG_PROMPT again.
```

```
get_generation_time(
 llm,
 sampling_params,
 LONG_PROMPT + "Question: what is the age of Zack Blue? Your answer: The age of Zack Blue
is ",
)
...
```

## ## Example workloads

We describe two example workloads, where APC can provide huge performance benefit:

- Long document query, where the user repeatedly queries the same long document (e.g. software manual or annual report) with different queries. In this case, instead of processing the long

document again and again, APC allows vLLM to process this long document \*only once\*, and all future requests can avoid recomputing this long document by reusing its KV cache. This allows vLLM to serve future requests with much higher throughput and much lower latency.

- Multi-round conversation, where the user may chat with the application multiple times in the same chatting session. In this case, instead of processing the whole chatting history again and again, APC allows vLLM to reuse the processing results of the chat history across all future rounds of conversation, allowing vLLM to serve future requests with much higher throughput and much lower latency.

## ## Limits

APC in general does not reduce the performance of vLLM. With that being said, APC only reduces the time of processing the queries (the prefilling phase) and does not reduce the time of generating new tokens (the decoding phase). So APC does not bring performance gain when vLLM spends most of the time generating answers to the queries (e.g. when the length of the answer is long), or new queries do not share the same prefix with any of existing queries (so that the computation cannot be reused).