[![Cornell University](https://static.arxiv.org/static/browse/0.3.4/images/icons/cu/cornell-reduced-white-SMALL.svg)](https://www.cornell.edu/)

We gratefully acknowledge support from the Simons Foundation, [member institutions](https://info.arxiv.org/about/ourmembers.html), and all contributors. [Donate](https://info.arxiv.org/about/donate.html)

[]({url_path\('ignore_me'\)})

[arXiv.org](/) > article trackbacks

[Help](https://info.arxiv.org/help) | [Advanced Search](https://arxiv.org/search/advanced)

All fields Title Author Abstract Comments Journal reference ACM classification MSC classification Report number arXiv identifier DOI ORCID arXiv author ID Help pages Full text

Search

[![arXiv logo](https://static.arxiv.org/static/browse/0.3.4/images/arxiv-logomark-small-white.svg)](https://arxiv.org/)

[![Cornell University Logo](https://static.arxiv.org/static/browse/0.3.4/images/icons/cu/cornell-

quick links

  * [Login](https://arxiv.org/login)

  * [Help Pages](https://info.arxiv.org/help)

  * [About](https://info.arxiv.org/about)

arXiv

Trackbacks

Trackbacks indicate external web sites that link to articles in arXiv.org. Trackbacks do not reflect the opinion of arXiv.org and may not reflect the opinions of that article's authors.

Trackback guide

By sending a [trackback](https://info.arxiv.org/help/trackback.html), you can notify arXiv.org that you have created a web page that references a paper. Popular blogging software supports trackback: you can send us a trackback about this paper by giving your software the following _trackback URL:_

> https://arxiv.org/trackback/{arXiv_id}

Some blogging software supports _trackback autodiscovery_ \-- in this case,
your software will automatically send a trackback as soon as your create a
link to our abstract page. See our [trackback help
page](https://info.arxiv.org/help/trackback.html) for more information.

Trackbacks for [2106.09685](/abs/2106.09685)

[Local LLM Fine-Tuning on Mac (M1 16GB)](/tb/redirect/1868221/66b125ffb)

[Towards Data Science - Medium@
[towardsdatascience.com/loca...](/tb/redirect/1868221/66b125ffb)] trackback
posted Thu, 1 Aug 2024 15:22:13 UTC

[Understanding Low Rank Adaptation (LoRA) in Fine Tuning
LLMs](/tb/redirect/1867703/1c8554eb7)

[Towards Data Science - Medium@
[towardsdatascience.com/unde...](/tb/redirect/1867703/1c8554eb7)] trackback
posted Fri, 24 May 2024 07:06:28 UTC

[LoRA: Revolutionizing Large Language Model Adaptation without Fine-
Tuning](/tb/redirect/1867348/8515a9574)

[Towards Data Science - Medium@

[towardsdatascience.com/lora...](/tb/redirect/1867348/8515a9574)] trackback

posted Tue, 23 Apr 2024 14:17:05 UTC

[An Overview of the LoRA Family](/tb/redirect/1866266/545a0f77c)

[Towards Data Science - Medium@

[towardsdatascience.com/an-o...](/tb/redirect/1866266/545a0f77c)] trackback

posted Sun, 10 Mar 2024 15:27:45 UTC

[Streamlining Giants](/tb/redirect/1866081/1df059787)

[Towards Data Science - Medium@

[towardsdatascience.com/stre...](/tb/redirect/1866081/1df059787)] trackback

posted Thu, 29 Feb 2024 08:20:22 UTC

[QLoRA -- How to Fine-Tune an LLM on a Single GPU](/tb/redirect/1865921/cf57b3a9a)

[Towards Data Science - Medium@

[towardsdatascience.com/qlor...](/tb/redirect/1865921/cf57b3a9a)] trackback

posted Thu, 22 Feb 2024 06:42:14 UTC

[LingoNaut Language Assistant](/tb/redirect/1865562/bf93c6a07)

[Towards Data Science - Medium@

[towardsdatascience.com/ling...](/tb/redirect/1865562/bf93c6a07)] trackback

posted Sun, 11 Feb 2024 15:17:09 UTC

[A Winding Road to Parameter Efficiency](/tb/redirect/1863792/43f01160d)

[Towards Data Science - Medium@

[towardsdatascience.com/a-wi...](/tb/redirect/1863792/43f01160d)] trackback

posted Thu, 4 Jan 2024 05:43:00 UTC

[Understanding LoRA -- Low Rank Adaptation For Finetuning Large

Models](/tb/redirect/1863369/d29a4d59c)

[Towards Data Science - Medium@

[towardsdatascience.com/unde...](/tb/redirect/1863369/d29a4d59c)] trackback

posted Fri, 22 Dec 2023 23:16:12 UTC

[Implementing LoRA from Scratch](/tb/redirect/1863049/83cf4537c)

[Towards Data Science - Medium@

[towardsdatascience.com/impl...](/tb/redirect/1863049/83cf4537c)] trackback

posted Tue, 12 Dec 2023 23:28:27 UTC

[The (local) unit of intelligence is

FLOPs](/tb/redirect/1849504/4d18fa4f9)

[Windows On Theory@

[windowsontheory.org/2023/06...](/tb/redirect/1849504/4d18fa4f9)] trackback

posted Mon, 5 Jun 2023 18:22:58 UTC

[QLoRa: Fine-Tune a Large Language Model on Your GPU](/tb/redirect/1849353/08d442f74)

[Towards Data Science - Medium@

[towardsdatascience.com/qlor...](/tb/redirect/1849353/08d442f74)] trackback

posted Tue, 30 May 2023 13:16:50 UTC

Click to view metadata for 2106.09685

> [Submitted on 17 Jun 2021 ([v1](https://arxiv.org/abs/2106.09685v1)), last

> revised 16 Oct 2021 (this version, v2)]

>

> # Title:LoRA: Low-Rank Adaptation of Large Language Models

>

> Authors:[Edward J.

> Hu](https://arxiv.org/search/cs?searchtype=author&query=Hu,+E+J), [Yelong

> Shen](https://arxiv.org/search/cs?searchtype=author&query=Shen,+Y), [Phillip

> Wallis](https://arxiv.org/search/cs?searchtype=author&query=Wallis,+P),

> [Zeyuan Allen-

> Zhu](https://arxiv.org/search/cs?searchtype=author&query=Allen-Zhu,+Z),

> [Yuanzhi Li](https://arxiv.org/search/cs?searchtype=author&query=Li,+Y),

> [Shean Wang](https://arxiv.org/search/cs?searchtype=author&query=Wang,+S),

> [Lu Wang](https://arxiv.org/search/cs?searchtype=author&query=Wang,+L),

> [Weizhu Chen](https://arxiv.org/search/cs?searchtype=author&query=Chen,+W)

>

>> Abstract:

>

> Comments: | Draft V2 includes better baselines, experiments on GLUE, and more on adapter latency

> ---|---

> Subjects: | Computation and Language (cs.CL); Artificial Intelligence (cs.AI); Machine Learning (cs.LG)

> Cite as: | [arXiv:2106.09685](https://arxiv.org/abs/2106.09685) [cs.CL]

> | (or [arXiv:2106.09685v2](https://arxiv.org/abs/2106.09685v2) [cs.CL] for this version)