

## NCCL

[2.25](https://docs.nvidia.com/deeplearning/sdk/nccl-archived/index.html)

- \* [Overview of NCCL](overview.html)

- \* [Setup](setup.html)

- \* [Using NCCL](usage.html)

- \* [Creating a Communicator](usage/communicators.html)

- \* [Creating a communicator with options](usage/communicators.html#creating-a-communicator-with-options)

- \* [Creating a communicator using multiple ncclUniqueIds](usage/communicators.html#creating-a-communicator-using-multiple-nccluniqueids)

- \* [Creating more communicators](usage/communicators.html#creating-more-communicators)

- \* [Using multiple NCCL communicators concurrently](usage/communicators.html#using-multiple-nccl-communicators-concurrently)

- \* [Finalizing a communicator](usage/communicators.html#finalizing-a-communicator)

- \* [Destroying a communicator](usage/communicators.html#destroying-a-communicator)

- \* [Error handling and communicator abort](usage/communicators.html#error-handling-and-communicator-abort)

- \* [Asynchronous errors and error handling](usage/communicators.html#asynchronous-errors-and-error-handling)

- \* [Fault Tolerance](usage/communicators.html#fault-tolerance)

- \* [Collective Operations](usage/collectives.html)

- \* [AllReduce](usage/collectives.html#allreduce)

- \* [Broadcast](usage/collectives.html#broadcast)

- \* [Reduce](usage/collectives.html#reduce)

- \* [AllGather](usage/collectives.html#allgather)

- \* [\[ReduceScatter\]\(usage/collectives.html#reducescatter\)](#)
- \* [\[Data Pointers\]\(usage/data.html\)](#)
- \* [\[CUDA Stream Semantics\]\(usage/streams.html\)](#)
  - \* [\[Mixing Multiple Streams within the same ncclGroupStart/End\(\) group\]\(usage/streams.html#mixing-multiple-streams-within-the-same-ncclgroupstart-end-group\)](#)
- \* [\[Group Calls\]\(usage/groups.html\)](#)
  - \* [\[Management Of Multiple GPUs From One Thread\]\(usage/groups.html#management-of-multiple-gpus-from-one-thread\)](#)
  - \* [\[Aggregated Operations \(2.2 and later\)\]\(usage/groups.html#aggregated-operations-2-2-and-later\)](#)
- \* [\[Nonblocking Group Operation\]\(usage/groups.html#nonblocking-group-operation\)](#)
- \* [\[Point-to-point communication\]\(usage/p2p.html\)](#)
- \* [\[Sendrecv\]\(usage/p2p.html#sendrecv\)](#)
- \* [\[One-to-all \(scatter\)\]\(usage/p2p.html#one-to-all-scatter\)](#)
- \* [\[All-to-one \(gather\)\]\(usage/p2p.html#all-to-one-gather\)](#)
- \* [\[All-to-all\]\(usage/p2p.html#all-to-all\)](#)
- \* [\[Neighbor exchange\]\(usage/p2p.html#neighbor-exchange\)](#)
- \* [\[Thread Safety\]\(usage/threadsafety.html\)](#)
- \* [\[In-place Operations\]\(usage/inplace.html\)](#)
- \* [\[Using NCCL with CUDA Graphs\]\(usage/cudagraph.html\)](#)
- \* [\[User Buffer Registration\]\(usage/bufferreg.html\)](#)
  - \* [\[NVLink Sharp Buffer Registration\]\(usage/bufferreg.html#nvlink-sharp-buffer-registration\)](#)
  - \* [\[IB Sharp Buffer Registration\]\(usage/bufferreg.html#ib-sharp-buffer-registration\)](#)
  - \* [\[General Buffer Registration\]\(usage/bufferreg.html#general-buffer-registration\)](#)
  - \* [\[Memory Allocator\]\(usage/bufferreg.html#memory-allocator\)](#)
- \* [\[NCCL API\]\(api.html\)](#)
  - \* [\[Communicator Creation and Management Functions\]\(api/comms.html\)](#)

- \* [\[ncclGetLastError\]\(api/comms.html#ncclgetlasterror\)](#)
- \* [\[ncclGetErrorString\]\(api/comms.html#ncclgeterrorstring\)](#)
- \* [\[ncclGetVersion\]\(api/comms.html#ncclgetversion\)](#)
- \* [\[ncclGetUniqueId\]\(api/comms.html#ncclgetuniqueid\)](#)
- \* [\[ncclCommInitRank\]\(api/comms.html#ncclcomminitrank\)](#)
- \* [\[ncclCommInitAll\]\(api/comms.html#ncclcomminitall\)](#)
- \* [\[ncclCommInitRankConfig\]\(api/comms.html#ncclcomminitrankconfig\)](#)
- \* [\[ncclCommInitRankScalable\]\(api/comms.html#ncclcomminitrankscalable\)](#)
- \* [\[ncclCommSplit\]\(api/comms.html#ncclcommsplit\)](#)
- \* [\[ncclCommFinalize\]\(api/comms.html#ncclcommfinalize\)](#)
- \* [\[ncclCommDestroy\]\(api/comms.html#ncclcommdestroy\)](#)
- \* [\[ncclCommAbort\]\(api/comms.html#ncclcommabort\)](#)
- \* [\[ncclCommGetAsyncError\]\(api/comms.html#ncclcommgetasynccerror\)](#)
- \* [\[ncclCommCount\]\(api/comms.html#ncclcommcount\)](#)
- \* [\[ncclCommCuDevice\]\(api/comms.html#ncclcommcudevice\)](#)
- \* [\[ncclCommUserRank\]\(api/comms.html#ncclcommuserrank\)](#)
- \* [\[ncclCommRegister\]\(api/comms.html#ncclcommregister\)](#)
- \* [\[ncclCommDeregister\]\(api/comms.html#ncclcommderegister\)](#)
- \* [\[ncclMemAlloc\]\(api/comms.html#ncclmemalloc\)](#)
- \* [\[ncclMemFree\]\(api/comms.html#ncclmemfree\)](#)
- \* [\[Collective Communication Functions\]\(api/colls.html\)](#)
  - \* [\[ncclAllReduce\]\(api/colls.html#ncclallreduce\)](#)
  - \* [\[ncclBroadcast\]\(api/colls.html#ncclbroadcast\)](#)
  - \* [\[ncclReduce\]\(api/colls.html#ncclreduce\)](#)
  - \* [\[ncclAllGather\]\(api/colls.html#ncclallgather\)](#)
  - \* [\[ncclReduceScatter\]\(api/colls.html#ncclreducescatter\)](#)
- \* [\[Group Calls\]\(api/group.html\)](#)

- \* [\[ncclGroupStart\]\(api/group.html#ncclgroupstart\)](#)
- \* [\[ncclGroupEnd\]\(api/group.html#ncclgroupend\)](#)
- \* [\[ncclGroupSimulateEnd\]\(api/group.html#ncclgroupsimulateend\)](#)
- \* [\[Point To Point Communication Functions\]\(api/p2p.html\)](#)
  - \* [\[ncclSend\]\(api/p2p.html#ncclsend\)](#)
  - \* [\[ncclRecv\]\(api/p2p.html#ncclrecv\)](#)
- \* [\[Types\]\(api/types.html\)](#)
  - \* [\[ncclComm\\_t\]\(api/types.html#ncclcomm-t\)](#)
  - \* [\[ncclResult\\_t\]\(api/types.html#ncclresult-t\)](#)
  - \* [\[ncclDataType\\_t\]\(api/types.html#nccldatatype-t\)](#)
  - \* [\[ncclRedOp\\_t\]\(api/types.html#ncclredop-t\)](#)
  - \* [\[ncclScalarResidence\\_t\]\(api/types.html#ncclscalarresidence-t\)](#)
  - \* [\[ncclConfig\\_t\]\(api/types.html#ncclconfig-t\)](#)
  - \* [\[ncclSimInfo\\_t\]\(api/types.html#ncclsiminfo-t\)](#)
- \* [\[User Defined Reduction Operators\]\(api/ops.html\)](#)
  - \* [\[ncclRedOpCreatePreMulSum\]\(api/ops.html#ncclredopcreatepremulsum\)](#)
  - \* [\[ncclRedOpDestroy\]\(api/ops.html#ncclredopdestroy\)](#)
- \* [\[Migrating from NCCL 1 to NCCL 2\]\(nccl1.html\)](#)
  - \* [\[Initialization\]\(nccl1.html#initialization\)](#)
  - \* [\[Communication\]\(nccl1.html#communication\)](#)
  - \* [\[Counts\]\(nccl1.html#counts\)](#)
  - \* [\[In-place usage for AllGather and ReduceScatter\]\(nccl1.html#in-place-usage-for-allgather-and-reducescatter\)](#)
  - \* [\[AllGather arguments order\]\(nccl1.html#allgather-arguments-order\)](#)
  - \* [\[Datatypes\]\(nccl1.html#datatypes\)](#)
  - \* [\[Error codes\]\(nccl1.html#error-codes\)](#)
  - \* [\[Examples\]\(examples.html\)](#)

\* [Communicator Creation and Destruction

Examples](examples.html#communicator-creation-and-destruction-examples)

\* [Example 1: Single Process, Single Thread, Multiple

Devices](examples.html#example-1-single-process-single-thread-multiple-devices)

\* [Example 2: One Device per Process or

Thread](examples.html#example-2-one-device-per-process-or-thread)

\* [Example 3: Multiple Devices per

Thread](examples.html#example-3-multiple-devices-per-thread)

\* [Example 4: Multiple communicators per

device](examples.html#example-4-multiple-communicators-per-device)

\* [Communication Examples](examples.html#communication-examples)

\* [Example 1: One Device per Process or

Thread](examples.html#example-1-one-device-per-process-or-thread)

\* [Example 2: Multiple Devices per

Thread](examples.html#example-2-multiple-devices-per-thread)

\* [NCCL and MPI](mpi.html)

\* [API](mpi.html#api)

\* [Using multiple devices per process](mpi.html#using-multiple-devices-per-process)

\* [ReduceScatter operation](mpi.html#reducescatter-operation)

\* [Send and Receive counts](mpi.html#send-and-receive-counts)

\* [Other collectives and point-to-point

operations](mpi.html#other-collectives-and-point-to-point-operations)

\* [In-place operations](mpi.html#in-place-operations)

\* [Using NCCL within an MPI Program](mpi.html#using-nccl-within-an-mpi-program)

\* [MPI Progress](mpi.html#mpi-progress)

\* [Inter-GPU Communication with CUDA-aware

MPI](mpi.html#inter-gpu-communication-with-cuda-aware-mpi)

- \* [Environment Variables](env.html)
- \* [System configuration](env.html#system-configuration)
- \* [NCCL\_SOCKET\_IFNAME](env.html#nccl-socket-ifname)
  - \* [Values accepted](env.html#values-accepted)
- \* [NCCL\_SOCKET\_FAMILY](env.html#nccl-socket-family)
  - \* [Values accepted](env.html#id2)
- \* [NCCL\_SOCKET\_RETRY\_CNT](env.html#nccl-socket-retry-cnt)
  - \* [Values accepted](env.html#id3)
- \* [NCCL\_SOCKET\_RETRY\_SLEEP\_MSEC](env.html#nccl-socket-retry-sleep-msec)
  - \* [Values accepted](env.html#id4)
- \* [NCCL\_SOCKET\_NTHREADS](env.html#nccl-socket-nthreads)
  - \* [Values accepted](env.html#id5)
- \* [NCCL\_NSOCKS\_PERTHREAD](env.html#nccl-nsocks-perthread)
  - \* [Values accepted](env.html#id6)
- \* [NCCL\_CROSS\_NIC](env.html#nccl-cross-nic)
  - \* [Values accepted](env.html#id7)
- \* [NCCL\_IB\_HCA](env.html#nccl-ib-hca)
  - \* [Values accepted](env.html#id8)
- \* [NCCL\_IB\_TIMEOUT](env.html#nccl-ib-timeout)
  - \* [Values accepted](env.html#id9)
- \* [NCCL\_IB\_RETRY\_CNT](env.html#nccl-ib-retry-cnt)
  - \* [Values accepted](env.html#id10)
- \* [NCCL\_IB\_GID\_INDEX](env.html#nccl-ib-gid-index)
  - \* [Values accepted](env.html#id11)
- \* [NCCL\_IB\_ADDR\_FAMILY](env.html#nccl-ib-addr-family)
  - \* [Values accepted](env.html#id12)
- \* [NCCL\_IB\_ADDR\_RANGE](env.html#nccl-ib-addr-range)

\* [Values accepted](env.html#id13)

\* [NCCL\_IB\_ROCE\_VERSION\_NUM](env.html#nccl-ib-roce-version-num)

\* [Values accepted](env.html#id14)

\* [NCCL\_IB\_SL](env.html#nccl-ib-sl)

\* [Values accepted](env.html#id15)

\* [NCCL\_IB\_TC](env.html#nccl-ib-tc)

\* [Values accepted](env.html#id16)

\* [NCCL\_IB\_FIFO\_TC](env.html#nccl-ib-fifo-tc)

\* [Values accepted](env.html#id17)

\* [NCCL\_IB\_RETURN\_ASYNC\_EVENTS](env.html#nccl-ib-return-async-events)

\* [Values accepted](env.html#id18)

\* [NCCL\_OOB\_NET\_ENABLE](env.html#nccl-oob-net-enable)

\* [Values accepted](env.html#id19)

\* [NCCL\_OOB\_NET\_IFNAME](env.html#nccl-oob-net-ifname)

\* [Values accepted](env.html#id20)

\* [NCCL\_UID\_STAGGER\_THRESHOLD](env.html#nccl-uid-stagger-threshold)

\* [Values accepted](env.html#id21)

\* [NCCL\_UID\_STAGGER\_RATE](env.html#nccl-uid-stagger-rate)

\* [Values accepted](env.html#id22)

\* [NCCL\_NET](env.html#nccl-net)

\* [Values accepted](env.html#id23)

\* [NCCL\_NET\_PLUGIN](env.html#nccl-net-plugin)

\* [Values accepted](env.html#id24)

\* [NCCL\_TUNER\_PLUGIN](env.html#nccl-tuner-plugin)

\* [Values accepted](env.html#id25)

\* [NCCL\_PROFILER\_PLUGIN](env.html#nccl-profiler-plugin)

\* [Values accepted](env.html#id26)

- \* [NCCL\_IGNORE\_CPU\_AFFINITY](env.html#nccl-ignore-cpu-affinity)
  - \* [Values accepted](env.html#id27)
- \* [NCCL\_CONF\_FILE](env.html#nccl-conf-file)
  - \* [Values accepted](env.html#id28)
- \* [NCCL\_DEBUG](env.html#nccl-debug)
  - \* [Values accepted](env.html#id30)
- \* [NCCL\_DEBUG\_FILE](env.html#nccl-debug-file)
  - \* [Values accepted](env.html#id31)
- \* [NCCL\_DEBUG\_SUBSYS](env.html#nccl-debug-subsys)
  - \* [Values accepted](env.html#id32)
- \* [NCCL\_COLLNET\_ENABLE](env.html#nccl-collnet-enable)
  - \* [Value accepted](env.html#value-accepted)
- \* [NCCL\_COLLNET\_NODE\_THRESHOLD](env.html#nccl-collnet-node-threshold)
  - \* [Value accepted](env.html#id33)
- \* [NCCL\_TOPO\_FILE](env.html#nccl-topo-file)
  - \* [Value accepted](env.html#id34)
- \* [NCCL\_TOPO\_DUMP\_FILE](env.html#nccl-topo-dump-file)
  - \* [Value accepted](env.html#id35)
- \* [NCCL\_SET\_THREAD\_NAME](env.html#nccl-set-thread-name)
  - \* [Value accepted](env.html#id36)
- \* [Debugging](env.html#debugging)
- \* [NCCL\_P2P\_DISABLE](env.html#nccl-p2p-disable)
  - \* [Values accepted](env.html#id37)
- \* [NCCL\_P2P\_LEVEL](env.html#nccl-p2p-level)
  - \* [Values accepted](env.html#id38)
  - \* [Integer Values (Legacy)](env.html#integer-values-legacy)
- \* [NCCL\_P2P\_DIRECT\_DISABLE](env.html#nccl-p2p-direct-disable)



\* [Values accepted](env.html#id39)

\* [NCCL\_SHM\_DISABLE](env.html#nccl-shm-disable)

\* [Values accepted](env.html#id40)

\* [NCCL\_BUFFSIZE](env.html#nccl-buffersize)

\* [Values accepted](env.html#id41)

\* [NCCL\_NTHREADS](env.html#nccl-nthreads)

\* [Values accepted](env.html#id42)

\* [NCCL\_MAX\_NCHANNELS](env.html#nccl-max-nchannels)

\* [Values accepted](env.html#id43)

\* [NCCL\_MIN\_NCHANNELS](env.html#nccl-min-nchannels)

\* [Values accepted](env.html#id44)

\* [NCCL\_CHECKS\_DISABLE](env.html#nccl-checks-disable)

\* [Values accepted](env.html#id45)

\* [NCCL\_CHECK\_POINTERS](env.html#nccl-check-pointers)

\* [Values accepted](env.html#id46)

\* [NCCL\_LAUNCH\_MODE](env.html#nccl-launch-mode)

\* [Values accepted](env.html#id47)

\* [NCCL\_IB\_DISABLE](env.html#nccl-ib-disable)

\* [Values accepted](env.html#id48)

\* [NCCL\_IB\_AR\_THRESHOLD](env.html#nccl-ib-ar-threshold)

\* [Values accepted](env.html#id49)

\* [NCCL\_IB\_QPS\_PER\_CONNECTION](env.html#nccl-ib-qps-per-connection)

\* [Values accepted](env.html#id50)

\* [NCCL\_IB\_SPLIT\_DATA\_ON\_QPS](env.html#nccl-ib-split-data-on-qps)

\* [Values accepted](env.html#id51)

\* [NCCL\_IB\_CUDA\_SUPPORT](env.html#nccl-ib-cuda-support)

\* [Values accepted](env.html#id52)

\* [NCCL\_IB\_PCI\_RELAXED\_ORDERING](env.html#nccl-ib-pci-relaxed-ordering)

\* [Values accepted](env.html#id53)

\* [NCCL\_IB\_ADAPTIVE\_ROUTING](env.html#nccl-ib-adaptive-routing)

\* [Values accepted](env.html#id54)

\* [NCCL\_IB\_ECE\_ENABLE](env.html#nccl-ib-ece-enable)

\* [Values accepted](env.html#id55)

\* [NCCL\_MEM\_SYNC\_DOMAIN](env.html#nccl-mem-sync-domain)

\* [Values accepted](env.html#id56)

\* [NCCL\_CUMEM\_ENABLE](env.html#nccl-cumem-enable)

\* [Values accepted](env.html#id57)

\* [NCCL\_CUMEM\_HOST\_ENABLE](env.html#nccl-cumem-host-enable)

\* [Values accepted](env.html#id58)

\* [NCCL\_NET\_GDR\_LEVEL (formerly

NCCL\_IB\_GDR\_LEVEL)](env.html#nccl-net-gdr-level-formerly-nccl-ib-gdr-level)

\* [Values accepted](env.html#id59)

\* [Integer Values (Legacy)](env.html#id60)

\* [NCCL\_NET\_GDR\_READ](env.html#nccl-net-gdr-read)

\* [Values accepted](env.html#id61)

\* [NCCL\_NET\_SHARED\_BUFFERS](env.html#nccl-net-shared-buffers)

\* [Value accepted](env.html#id62)

\* [NCCL\_NET\_SHARED\_COMMS](env.html#nccl-net-shared-comms)

\* [Value accepted](env.html#id63)

\* [NCCL\_SINGLE\_RING\_THRESHOLD](env.html#nccl-single-ring-threshold)

\* [Values accepted](env.html#id64)

\* [NCCL\_LL\_THRESHOLD](env.html#nccl-ll-threshold)

\* [Values accepted](env.html#id65)

\* [NCCL\_TREE\_THRESHOLD](env.html#nccl-tree-threshold)

\* [Values accepted](env.html#id66)

\* [NCCL\_ALGO](env.html#nccl-algo)

\* [Values accepted](env.html#id67)

\* [NCCL\_PROTO](env.html#nccl-proto)

\* [Values accepted](env.html#id68)

\* [NCCL\_NVB\_DISABLE](env.html#nccl-nvb-disable)

\* [Value accepted](env.html#id69)

\* [NCCL\_PXN\_DISABLE](env.html#nccl-pxn-disable)

\* [Value accepted](env.html#id70)

\* [NCCL\_P2P\_PXN\_LEVEL](env.html#nccl-p2p-pxn-level)

\* [Value accepted](env.html#id71)

\* [NCCL\_RUNTIME\_CONNECT](env.html#nccl-runtime-connect)

\* [Value accepted](env.html#id72)

\* [NCCL\_GRAPH\_REGISTER](env.html#nccl-graph-register)

\* [Value accepted](env.html#id74)

\* [NCCL\_LOCAL\_REGISTER](env.html#nccl-local-register)

\* [Value accepted](env.html#id75)

\* [NCCL\_LEGACY\_CUDA\_REGISTER](env.html#nccl-legacy-cuda-register)

\* [Value accepted](env.html#id76)

\* [NCCL\_SET\_STACK\_SIZE](env.html#nccl-set-stack-size)

\* [Value accepted](env.html#id77)

\* [NCCL\_GRAPH\_MIXING\_SUPPORT](env.html#nccl-graph-mixing-support)

\* [Value accepted](env.html#id79)

\* [NCCL\_DMABUF\_ENABLE](env.html#nccl-dmabuf-enable)

\* [Value accepted](env.html#id80)

\* [NCCL\_P2P\_NET\_CHUNKSIZE](env.html#nccl-p2p-net-chunksize)

\* [Values accepted](env.html#id81)

- \* [\[NCCL\\_P2P\\_LL\\_THRESHOLD\]\(env.html#nccl-p2p-ll-threshold\)](#)
  - \* [\[Values accepted\]\(env.html#id82\)](#)
- \* [\[NCCL\\_ALLOC\\_P2P\\_NET\\_LL\\_BUFFERS\]\(env.html#nccl-alloc-p2p-net-ll-buffers\)](#)
  - \* [\[Values accepted\]\(env.html#id83\)](#)
- \* [\[NCCL\\_COMM\\_BLOCKING\]\(env.html#nccl-comm-blocking\)](#)
  - \* [\[Values accepted\]\(env.html#id84\)](#)
- \* [\[NCCL\\_CGA\\_CLUSTER\\_SIZE\]\(env.html#nccl-cga-cluster-size\)](#)
  - \* [\[Values accepted\]\(env.html#id85\)](#)
- \* [\[NCCL\\_MAX\\_CTAS\]\(env.html#nccl-max-ctas\)](#)
  - \* [\[Values accepted\]\(env.html#id86\)](#)
- \* [\[NCCL\\_MIN\\_CTAS\]\(env.html#nccl-min-ctas\)](#)
  - \* [\[Values accepted\]\(env.html#id87\)](#)
- \* [\[NCCL\\_NVLS\\_ENABLE\]\(env.html#nccl-nvls-enable\)](#)
  - \* [\[Values accepted\]\(env.html#id88\)](#)
- \* [\[NCCL\\_IB\\_MERGE\\_NICS\]\(env.html#nccl-ib-merge-nics\)](#)
  - \* [\[Values accepted\]\(env.html#id89\)](#)
- \* [\[NCCL\\_MNNVL\\_ENABLE\]\(env.html#nccl-mnnvl-enable\)](#)
  - \* [\[Values accepted\]\(env.html#id90\)](#)
- \* [\[NCCL\\_RAS\\_ENABLE\]\(env.html#nccl-ras-enable\)](#)
  - \* [\[Values accepted\]\(env.html#id91\)](#)
- \* [\[NCCL\\_RAS\\_ADDR\]\(env.html#nccl-ras-addr\)](#)
  - \* [\[Values accepted\]\(env.html#id92\)](#)
- \* [\[NCCL\\_RAS\\_TIMEOUT\\_FACTOR\]\(env.html#nccl-ras-timeout-factor\)](#)
  - \* [\[Values accepted\]\(env.html#id93\)](#)
- \* [\[Troubleshooting\]\(troubleshooting.html\)](#)
  - \* [\[Errors\]\(troubleshooting.html#errors\)](#)
  - \* [\[RAS\]\(troubleshooting.html#ras\)](#)

- \* [\[RAS\]\(troubleshooting/ras.html\)](#)
- \* [\[Principle of Operation\]\(troubleshooting/ras.html#principle-of-operation\)](#)
- \* [\[RAS Queries\]\(troubleshooting/ras.html#ras-queries\)](#)
- \* [\[Sample Output\]\(troubleshooting/ras.html#sample-output\)](#)
- \* [\[GPU Direct\]\(troubleshooting.html#gpu-direct\)](#)
- \* [\[GPU-to-GPU communication\]\(troubleshooting.html#gpu-to-gpu-communication\)](#)
- \* [\[GPU-to-NIC communication\]\(troubleshooting.html#gpu-to-nic-communication\)](#)
- \* [\[PCI Access Control Services \(ACS\)\]\(troubleshooting.html#pci-access-control-services-ac\)](#)
- \* [\[Topology detection\]\(troubleshooting.html#topology-detection\)](#)
- \* [\[Shared memory\]\(troubleshooting.html#shared-memory\)](#)
- \* [\[Docker\]\(troubleshooting.html#docker\)](#)
- \* [\[Systemd\]\(troubleshooting.html#systemd\)](#)
- \* [\[Networking issues\]\(troubleshooting.html#networking-issues\)](#)
- \* [\[IP Network Interfaces\]\(troubleshooting.html#ip-network-interfaces\)](#)
- \* [\[IP Ports\]\(troubleshooting.html#ip-ports\)](#)
- \* [\[InfiniBand\]\(troubleshooting.html#infiniband\)](#)
- \* [\[RDMA over Converged Ethernet \(RoCE\)\]\(troubleshooting.html#rdma-over-converged-ethernet-roce\)](#)

\_\_NCCL

- \* [Docs »](#)
- \* [NVIDIA Collective Communication Library \(NCCL\) Documentation](#)
- \* [\[ View page source\]\(\\_sources/index.rst.txt\)](#)

\* \* \*

## Contents:

- \* [\[Overview of NCCL\]\(overview.html\)](#)
- \* [\[Setup\]\(setup.html\)](#)
- \* [\[Using NCCL\]\(usage.html\)](#)
- \* [\[Creating a Communicator\]\(usage/communicators.html\)](#)
  - \* [\[Creating a communicator with options\]\(usage/communicators.html#creating-a-communicator-with-options\)](#)
  - \* [\[Creating a communicator using multiple ncclUniqeIds\]\(usage/communicators.html#creating-a-communicator-using-multiple-nccluniqueids\)](#)
  - \* [\[Creating more communicators\]\(usage/communicators.html#creating-more-communicators\)](#)
  - \* [\[Using multiple NCCL communicators concurrently\]\(usage/communicators.html#using-multiple-nccl-communicators-concurrently\)](#)
  - \* [\[Finalizing a communicator\]\(usage/communicators.html#finalizing-a-communicator\)](#)
  - \* [\[Destroying a communicator\]\(usage/communicators.html#destroying-a-communicator\)](#)
  - \* [\[Error handling and communicator abort\]\(usage/communicators.html#error-handling-and-communicator-abort\)](#)
  - \* [\[Asynchronous errors and error handling\]\(usage/communicators.html#asynchronous-errors-and-error-handling\)](#)
  - \* [\[Fault Tolerance\]\(usage/communicators.html#fault-tolerance\)](#)
- \* [\[Collective Operations\]\(usage/collectives.html\)](#)
  - \* [\[AllReduce\]\(usage/collectives.html#allreduce\)](#)
  - \* [\[Broadcast\]\(usage/collectives.html#broadcast\)](#)
  - \* [\[Reduce\]\(usage/collectives.html#reduce\)](#)
  - \* [\[AllGather\]\(usage/collectives.html#allgather\)](#)

- \* [\[ReduceScatter\]\(usage/collectives.html#reducescatter\)](#)
- \* [\[Data Pointers\]\(usage/data.html\)](#)
- \* [\[CUDA Stream Semantics\]\(usage/streams.html\)](#)
  - \* [\[Mixing Multiple Streams within the same ncclGroupStart/End\(\) group\]\(usage/streams.html#mixing-multiple-streams-within-the-same-ncclgroupstart-end-group\)](#)
- \* [\[Group Calls\]\(usage/groups.html\)](#)
  - \* [\[Management Of Multiple GPUs From One Thread\]\(usage/groups.html#management-of-multiple-gpus-from-one-thread\)](#)
  - \* [\[Aggregated Operations \(2.2 and later\)\]\(usage/groups.html#aggregated-operations-2-2-and-later\)](#)
- \* [\[Nonblocking Group Operation\]\(usage/groups.html#nonblocking-group-operation\)](#)
- \* [\[Point-to-point communication\]\(usage/p2p.html\)](#)
- \* [\[Sendrecv\]\(usage/p2p.html#sendrecv\)](#)
- \* [\[One-to-all \(scatter\)\]\(usage/p2p.html#one-to-all-scatter\)](#)
- \* [\[All-to-one \(gather\)\]\(usage/p2p.html#all-to-one-gather\)](#)
- \* [\[All-to-all\]\(usage/p2p.html#all-to-all\)](#)
- \* [\[Neighbor exchange\]\(usage/p2p.html#neighbor-exchange\)](#)
- \* [\[Thread Safety\]\(usage/threadsafety.html\)](#)
- \* [\[In-place Operations\]\(usage/inplace.html\)](#)
- \* [\[Using NCCL with CUDA Graphs\]\(usage/cudagraph.html\)](#)
- \* [\[User Buffer Registration\]\(usage/bufferreg.html\)](#)
  - \* [\[NVLink Sharp Buffer Registration\]\(usage/bufferreg.html#nvlink-sharp-buffer-registration\)](#)
  - \* [\[IB Sharp Buffer Registration\]\(usage/bufferreg.html#ib-sharp-buffer-registration\)](#)
  - \* [\[General Buffer Registration\]\(usage/bufferreg.html#general-buffer-registration\)](#)
  - \* [\[Memory Allocator\]\(usage/bufferreg.html#memory-allocator\)](#)
- \* [\[NCCL API\]\(api.html\)](#)
  - \* [\[Communicator Creation and Management Functions\]\(api/comms.html\)](#)

- \* [\[ncclGetLastError\]\(api/comms.html#ncclgetlasterror\)](#)
- \* [\[ncclGetErrorString\]\(api/comms.html#ncclgeterrorstring\)](#)
- \* [\[ncclGetVersion\]\(api/comms.html#ncclgetversion\)](#)
- \* [\[ncclGetUniqueId\]\(api/comms.html#ncclgetuniqueid\)](#)
- \* [\[ncclCommInitRank\]\(api/comms.html#ncclcomminitrank\)](#)
- \* [\[ncclCommInitAll\]\(api/comms.html#ncclcomminitall\)](#)
- \* [\[ncclCommInitRankConfig\]\(api/comms.html#ncclcomminitrankconfig\)](#)
- \* [\[ncclCommInitRankScalable\]\(api/comms.html#ncclcomminitrankscalable\)](#)
- \* [\[ncclCommSplit\]\(api/comms.html#ncclcommsplit\)](#)
- \* [\[ncclCommFinalize\]\(api/comms.html#ncclcommfinalize\)](#)
- \* [\[ncclCommDestroy\]\(api/comms.html#ncclcommdestroy\)](#)
- \* [\[ncclCommAbort\]\(api/comms.html#ncclcommabort\)](#)
- \* [\[ncclCommGetAsyncError\]\(api/comms.html#ncclcommgetasynccerror\)](#)
- \* [\[ncclCommCount\]\(api/comms.html#ncclcommcount\)](#)
- \* [\[ncclCommCuDevice\]\(api/comms.html#ncclcommcudevice\)](#)
- \* [\[ncclCommUserRank\]\(api/comms.html#ncclcommuserrank\)](#)
- \* [\[ncclCommRegister\]\(api/comms.html#ncclcommregister\)](#)
- \* [\[ncclCommDeregister\]\(api/comms.html#ncclcommderegister\)](#)
- \* [\[ncclMemAlloc\]\(api/comms.html#ncclmemalloc\)](#)
- \* [\[ncclMemFree\]\(api/comms.html#ncclmemfree\)](#)
- \* [\[Collective Communication Functions\]\(api/colls.html\)](#)
  - \* [\[ncclAllReduce\]\(api/colls.html#ncclallreduce\)](#)
  - \* [\[ncclBroadcast\]\(api/colls.html#ncclbroadcast\)](#)
  - \* [\[ncclReduce\]\(api/colls.html#ncclreduce\)](#)
  - \* [\[ncclAllGather\]\(api/colls.html#ncclallgather\)](#)
  - \* [\[ncclReduceScatter\]\(api/colls.html#ncclreducescatter\)](#)
- \* [\[Group Calls\]\(api/group.html\)](#)



- \* [\[ncclGroupStart\]\(api/group.html#ncclgroupstart\)](#)
- \* [\[ncclGroupEnd\]\(api/group.html#ncclgroupend\)](#)
- \* [\[ncclGroupSimulateEnd\]\(api/group.html#ncclgroupsimulateend\)](#)
- \* [\[Point To Point Communication Functions\]\(api/p2p.html\)](#)
  - \* [\[ncclSend\]\(api/p2p.html#ncclsend\)](#)
  - \* [\[ncclRecv\]\(api/p2p.html#ncclrecv\)](#)
- \* [\[Types\]\(api/types.html\)](#)
  - \* [\[ncclComm\\_t\]\(api/types.html#ncclcomm-t\)](#)
  - \* [\[ncclResult\\_t\]\(api/types.html#ncclresult-t\)](#)
  - \* [\[ncclDataType\\_t\]\(api/types.html#nccldatatype-t\)](#)
  - \* [\[ncclRedOp\\_t\]\(api/types.html#ncclredop-t\)](#)
  - \* [\[ncclScalarResidence\\_t\]\(api/types.html#ncclscalarresidence-t\)](#)
  - \* [\[ncclConfig\\_t\]\(api/types.html#ncclconfig-t\)](#)
  - \* [\[ncclSimInfo\\_t\]\(api/types.html#ncclsiminfo-t\)](#)
- \* [\[User Defined Reduction Operators\]\(api/ops.html\)](#)
  - \* [\[ncclRedOpCreatePreMulSum\]\(api/ops.html#ncclredopcreatepremulsum\)](#)
  - \* [\[ncclRedOpDestroy\]\(api/ops.html#ncclredopdestroy\)](#)
- \* [\[Migrating from NCCL 1 to NCCL 2\]\(nccl1.html\)](#)
  - \* [\[Initialization\]\(nccl1.html#initialization\)](#)
  - \* [\[Communication\]\(nccl1.html#communication\)](#)
  - \* [\[Counts\]\(nccl1.html#counts\)](#)
  - \* [\[In-place usage for AllGather and ReduceScatter\]\(nccl1.html#in-place-usage-for-allgather-and-reducescatter\)](#)
  - \* [\[AllGather arguments order\]\(nccl1.html#allgather-arguments-order\)](#)
  - \* [\[Datatypes\]\(nccl1.html#datatypes\)](#)
  - \* [\[Error codes\]\(nccl1.html#error-codes\)](#)
  - \* [\[Examples\]\(examples.html\)](#)

\* [Communicator Creation and Destruction

Examples](examples.html#communicator-creation-and-destruction-examples)

\* [Example 1: Single Process, Single Thread, Multiple

Devices](examples.html#example-1-single-process-single-thread-multiple-devices)

\* [Example 2: One Device per Process or

Thread](examples.html#example-2-one-device-per-process-or-thread)

\* [Example 3: Multiple Devices per

Thread](examples.html#example-3-multiple-devices-per-thread)

\* [Example 4: Multiple communicators per

device](examples.html#example-4-multiple-communicators-per-device)

\* [Communication Examples](examples.html#communication-examples)

\* [Example 1: One Device per Process or

Thread](examples.html#example-1-one-device-per-process-or-thread)

\* [Example 2: Multiple Devices per

Thread](examples.html#example-2-multiple-devices-per-thread)

\* [NCCL and MPI](mpi.html)

\* [API](mpi.html#api)

\* [Using multiple devices per process](mpi.html#using-multiple-devices-per-process)

\* [ReduceScatter operation](mpi.html#reducescatter-operation)

\* [Send and Receive counts](mpi.html#send-and-receive-counts)

\* [Other collectives and point-to-point

operations](mpi.html#other-collectives-and-point-to-point-operations)

\* [In-place operations](mpi.html#in-place-operations)

\* [Using NCCL within an MPI Program](mpi.html#using-nccl-within-an-mpi-program)

\* [MPI Progress](mpi.html#mpi-progress)

\* [Inter-GPU Communication with CUDA-aware

MPI](mpi.html#inter-gpu-communication-with-cuda-aware-mpi)

- \* [Environment Variables](env.html)
- \* [System configuration](env.html#system-configuration)
- \* [NCCL\_SOCKET\_IFNAME](env.html#nccl-socket-ifname)
  - \* [Values accepted](env.html#values-accepted)
- \* [NCCL\_SOCKET\_FAMILY](env.html#nccl-socket-family)
  - \* [Values accepted](env.html#id2)
- \* [NCCL\_SOCKET\_RETRY\_CNT](env.html#nccl-socket-retry-cnt)
  - \* [Values accepted](env.html#id3)
- \* [NCCL\_SOCKET\_RETRY\_SLEEP\_MSEC](env.html#nccl-socket-retry-sleep-msec)
  - \* [Values accepted](env.html#id4)
- \* [NCCL\_SOCKET\_NTHREADS](env.html#nccl-socket-nthreads)
  - \* [Values accepted](env.html#id5)
- \* [NCCL\_NSOCKS\_PERTHREAD](env.html#nccl-nsocks-perthread)
  - \* [Values accepted](env.html#id6)
- \* [NCCL\_CROSS\_NIC](env.html#nccl-cross-nic)
  - \* [Values accepted](env.html#id7)
- \* [NCCL\_IB\_HCA](env.html#nccl-ib-hca)
  - \* [Values accepted](env.html#id8)
- \* [NCCL\_IB\_TIMEOUT](env.html#nccl-ib-timeout)
  - \* [Values accepted](env.html#id9)
- \* [NCCL\_IB\_RETRY\_CNT](env.html#nccl-ib-retry-cnt)
  - \* [Values accepted](env.html#id10)
- \* [NCCL\_IB\_GID\_INDEX](env.html#nccl-ib-gid-index)
  - \* [Values accepted](env.html#id11)
- \* [NCCL\_IB\_ADDR\_FAMILY](env.html#nccl-ib-addr-family)
  - \* [Values accepted](env.html#id12)
- \* [NCCL\_IB\_ADDR\_RANGE](env.html#nccl-ib-addr-range)

\* [Values accepted](env.html#id13)

\* [NCCL\_IB\_ROCE\_VERSION\_NUM](env.html#nccl-ib-roce-version-num)

\* [Values accepted](env.html#id14)

\* [NCCL\_IB\_SL](env.html#nccl-ib-sl)

\* [Values accepted](env.html#id15)

\* [NCCL\_IB\_TC](env.html#nccl-ib-tc)

\* [Values accepted](env.html#id16)

\* [NCCL\_IB\_FIFO\_TC](env.html#nccl-ib-fifo-tc)

\* [Values accepted](env.html#id17)

\* [NCCL\_IB\_RETURN\_ASYNC\_EVENTS](env.html#nccl-ib-return-async-events)

\* [Values accepted](env.html#id18)

\* [NCCL\_OOB\_NET\_ENABLE](env.html#nccl-oob-net-enable)

\* [Values accepted](env.html#id19)

\* [NCCL\_OOB\_NET\_IFNAME](env.html#nccl-oob-net-ifname)

\* [Values accepted](env.html#id20)

\* [NCCL\_UID\_STAGGER\_THRESHOLD](env.html#nccl-uid-stagger-threshold)

\* [Values accepted](env.html#id21)

\* [NCCL\_UID\_STAGGER\_RATE](env.html#nccl-uid-stagger-rate)

\* [Values accepted](env.html#id22)

\* [NCCL\_NET](env.html#nccl-net)

\* [Values accepted](env.html#id23)

\* [NCCL\_NET\_PLUGIN](env.html#nccl-net-plugin)

\* [Values accepted](env.html#id24)

\* [NCCL\_TUNER\_PLUGIN](env.html#nccl-tuner-plugin)

\* [Values accepted](env.html#id25)

\* [NCCL\_PROFILER\_PLUGIN](env.html#nccl-profiler-plugin)

\* [Values accepted](env.html#id26)

- \* [NCCL\_IGNORE\_CPU\_AFFINITY](env.html#nccl-ignore-cpu-affinity)
  - \* [Values accepted](env.html#id27)
- \* [NCCL\_CONF\_FILE](env.html#nccl-conf-file)
  - \* [Values accepted](env.html#id28)
- \* [NCCL\_DEBUG](env.html#nccl-debug)
  - \* [Values accepted](env.html#id30)
- \* [NCCL\_DEBUG\_FILE](env.html#nccl-debug-file)
  - \* [Values accepted](env.html#id31)
- \* [NCCL\_DEBUG\_SUBSYS](env.html#nccl-debug-subsys)
  - \* [Values accepted](env.html#id32)
- \* [NCCL\_COLLNET\_ENABLE](env.html#nccl-collnet-enable)
  - \* [Value accepted](env.html#value-accepted)
- \* [NCCL\_COLLNET\_NODE\_THRESHOLD](env.html#nccl-collnet-node-threshold)
  - \* [Value accepted](env.html#id33)
- \* [NCCL\_TOPO\_FILE](env.html#nccl-topo-file)
  - \* [Value accepted](env.html#id34)
- \* [NCCL\_TOPO\_DUMP\_FILE](env.html#nccl-topo-dump-file)
  - \* [Value accepted](env.html#id35)
- \* [NCCL\_SET\_THREAD\_NAME](env.html#nccl-set-thread-name)
  - \* [Value accepted](env.html#id36)
- \* [Debugging](env.html#debugging)
- \* [NCCL\_P2P\_DISABLE](env.html#nccl-p2p-disable)
  - \* [Values accepted](env.html#id37)
- \* [NCCL\_P2P\_LEVEL](env.html#nccl-p2p-level)
  - \* [Values accepted](env.html#id38)
  - \* [Integer Values (Legacy)](env.html#integer-values-legacy)
- \* [NCCL\_P2P\_DIRECT\_DISABLE](env.html#nccl-p2p-direct-disable)

\* [Values accepted](env.html#id39)

\* [NCCL\_SHM\_DISABLE](env.html#nccl-shm-disable)

\* [Values accepted](env.html#id40)

\* [NCCL\_BUFFSIZE](env.html#nccl-buffersize)

\* [Values accepted](env.html#id41)

\* [NCCL\_NTHREADS](env.html#nccl-nthreads)

\* [Values accepted](env.html#id42)

\* [NCCL\_MAX\_NCHANNELS](env.html#nccl-max-nchannels)

\* [Values accepted](env.html#id43)

\* [NCCL\_MIN\_NCHANNELS](env.html#nccl-min-nchannels)

\* [Values accepted](env.html#id44)

\* [NCCL\_CHECKS\_DISABLE](env.html#nccl-checks-disable)

\* [Values accepted](env.html#id45)

\* [NCCL\_CHECK\_POINTERS](env.html#nccl-check-pointers)

\* [Values accepted](env.html#id46)

\* [NCCL\_LAUNCH\_MODE](env.html#nccl-launch-mode)

\* [Values accepted](env.html#id47)

\* [NCCL\_IB\_DISABLE](env.html#nccl-ib-disable)

\* [Values accepted](env.html#id48)

\* [NCCL\_IB\_AR\_THRESHOLD](env.html#nccl-ib-ar-threshold)

\* [Values accepted](env.html#id49)

\* [NCCL\_IB\_QPS\_PER\_CONNECTION](env.html#nccl-ib-qps-per-connection)

\* [Values accepted](env.html#id50)

\* [NCCL\_IB\_SPLIT\_DATA\_ON\_QPS](env.html#nccl-ib-split-data-on-qps)

\* [Values accepted](env.html#id51)

\* [NCCL\_IB\_CUDA\_SUPPORT](env.html#nccl-ib-cuda-support)

\* [Values accepted](env.html#id52)

\* [NCCL\_IB\_PCI\_RELAXED\_ORDERING](env.html#nccl-ib-pci-relaxed-ordering)

\* [Values accepted](env.html#id53)

\* [NCCL\_IB\_ADAPTIVE\_ROUTING](env.html#nccl-ib-adaptive-routing)

\* [Values accepted](env.html#id54)

\* [NCCL\_IB\_ECE\_ENABLE](env.html#nccl-ib-ece-enable)

\* [Values accepted](env.html#id55)

\* [NCCL\_MEM\_SYNC\_DOMAIN](env.html#nccl-mem-sync-domain)

\* [Values accepted](env.html#id56)

\* [NCCL\_CUMEM\_ENABLE](env.html#nccl-cumem-enable)

\* [Values accepted](env.html#id57)

\* [NCCL\_CUMEM\_HOST\_ENABLE](env.html#nccl-cumem-host-enable)

\* [Values accepted](env.html#id58)

\* [NCCL\_NET\_GDR\_LEVEL] (formerly

NCCL\_IB\_GDR\_LEVEL))(env.html#nccl-net-gdr-level-formerly-nccl-ib-gdr-level)

\* [Values accepted](env.html#id59)

\* [Integer Values (Legacy)](env.html#id60)

\* [NCCL\_NET\_GDR\_READ](env.html#nccl-net-gdr-read)

\* [Values accepted](env.html#id61)

\* [NCCL\_NET\_SHARED\_BUFFERS](env.html#nccl-net-shared-buffers)

\* [Value accepted](env.html#id62)

\* [NCCL\_NET\_SHARED\_COMMS](env.html#nccl-net-shared-comms)

\* [Value accepted](env.html#id63)

\* [NCCL\_SINGLE\_RING\_THRESHOLD](env.html#nccl-single-ring-threshold)

\* [Values accepted](env.html#id64)

\* [NCCL\_LL\_THRESHOLD](env.html#nccl-ll-threshold)

\* [Values accepted](env.html#id65)

\* [NCCL\_TREE\_THRESHOLD](env.html#nccl-tree-threshold)

\* [Values accepted](env.html#id66)

\* [NCCL\_ALGO](env.html#nccl-algo)

\* [Values accepted](env.html#id67)

\* [NCCL\_PROTO](env.html#nccl-proto)

\* [Values accepted](env.html#id68)

\* [NCCL\_NVB\_DISABLE](env.html#nccl-nvb-disable)

\* [Value accepted](env.html#id69)

\* [NCCL\_PXN\_DISABLE](env.html#nccl-pxn-disable)

\* [Value accepted](env.html#id70)

\* [NCCL\_P2P\_PXN\_LEVEL](env.html#nccl-p2p-pxn-level)

\* [Value accepted](env.html#id71)

\* [NCCL\_RUNTIME\_CONNECT](env.html#nccl-runtime-connect)

\* [Value accepted](env.html#id72)

\* [NCCL\_GRAPH\_REGISTER](env.html#nccl-graph-register)

\* [Value accepted](env.html#id74)

\* [NCCL\_LOCAL\_REGISTER](env.html#nccl-local-register)

\* [Value accepted](env.html#id75)

\* [NCCL\_LEGACY\_CUDA\_REGISTER](env.html#nccl-legacy-cuda-register)

\* [Value accepted](env.html#id76)

\* [NCCL\_SET\_STACK\_SIZE](env.html#nccl-set-stack-size)

\* [Value accepted](env.html#id77)

\* [NCCL\_GRAPH\_MIXING\_SUPPORT](env.html#nccl-graph-mixing-support)

\* [Value accepted](env.html#id79)

\* [NCCL\_DMABUF\_ENABLE](env.html#nccl-dmabuf-enable)

\* [Value accepted](env.html#id80)

\* [NCCL\_P2P\_NET\_CHUNKSIZE](env.html#nccl-p2p-net-chunksize)

\* [Values accepted](env.html#id81)



- \* [\[NCCL\\_P2P\\_LL\\_THRESHOLD\]\(env.html#nccl-p2p-ll-threshold\)](#)
  - \* [\[Values accepted\]\(env.html#id82\)](#)
- \* [\[NCCL\\_ALLOC\\_P2P\\_NET\\_LL\\_BUFFERS\]\(env.html#nccl-alloc-p2p-net-ll-buffers\)](#)
  - \* [\[Values accepted\]\(env.html#id83\)](#)
- \* [\[NCCL\\_COMM\\_BLOCKING\]\(env.html#nccl-comm-blocking\)](#)
  - \* [\[Values accepted\]\(env.html#id84\)](#)
- \* [\[NCCL\\_CGA\\_CLUSTER\\_SIZE\]\(env.html#nccl-cga-cluster-size\)](#)
  - \* [\[Values accepted\]\(env.html#id85\)](#)
- \* [\[NCCL\\_MAX\\_CTAS\]\(env.html#nccl-max-ctas\)](#)
  - \* [\[Values accepted\]\(env.html#id86\)](#)
- \* [\[NCCL\\_MIN\\_CTAS\]\(env.html#nccl-min-ctas\)](#)
  - \* [\[Values accepted\]\(env.html#id87\)](#)
- \* [\[NCCL\\_NVLS\\_ENABLE\]\(env.html#nccl-nvls-enable\)](#)
  - \* [\[Values accepted\]\(env.html#id88\)](#)
- \* [\[NCCL\\_IB\\_MERGE\\_NICS\]\(env.html#nccl-ib-merge-nics\)](#)
  - \* [\[Values accepted\]\(env.html#id89\)](#)
- \* [\[NCCL\\_MNNVL\\_ENABLE\]\(env.html#nccl-mnnvl-enable\)](#)
  - \* [\[Values accepted\]\(env.html#id90\)](#)
- \* [\[NCCL\\_RAS\\_ENABLE\]\(env.html#nccl-ras-enable\)](#)
  - \* [\[Values accepted\]\(env.html#id91\)](#)
- \* [\[NCCL\\_RAS\\_ADDR\]\(env.html#nccl-ras-addr\)](#)
  - \* [\[Values accepted\]\(env.html#id92\)](#)
- \* [\[NCCL\\_RAS\\_TIMEOUT\\_FACTOR\]\(env.html#nccl-ras-timeout-factor\)](#)
  - \* [\[Values accepted\]\(env.html#id93\)](#)
- \* [\[Troubleshooting\]\(troubleshooting.html\)](#)
  - \* [\[Errors\]\(troubleshooting.html#errors\)](#)
  - \* [\[RAS\]\(troubleshooting.html#ras\)](#)

- \* [\[RAS\]\(troubleshooting/ras.html\)](#)
- \* [\[Principle of Operation\]\(troubleshooting/ras.html#principle-of-operation\)](#)
- \* [\[RAS Queries\]\(troubleshooting/ras.html#ras-queries\)](#)
- \* [\[Sample Output\]\(troubleshooting/ras.html#sample-output\)](#)
- \* [\[GPU Direct\]\(troubleshooting.html#gpu-direct\)](#)
- \* [\[GPU-to-GPU communication\]\(troubleshooting.html#gpu-to-gpu-communication\)](#)
- \* [\[GPU-to-NIC communication\]\(troubleshooting.html#gpu-to-nic-communication\)](#)
- \* [\[PCI Access Control Services \(ACS\)\]\(troubleshooting.html#pci-access-control-services-ac\)](#)
- \* [\[Topology detection\]\(troubleshooting.html#topology-detection\)](#)
- \* [\[Shared memory\]\(troubleshooting.html#shared-memory\)](#)
- \* [\[Docker\]\(troubleshooting.html#docker\)](#)
- \* [\[Systemd\]\(troubleshooting.html#systemd\)](#)
- \* [\[Networking issues\]\(troubleshooting.html#networking-issues\)](#)
- \* [\[IP Network Interfaces\]\(troubleshooting.html#ip-network-interfaces\)](#)
- \* [\[IP Ports\]\(troubleshooting.html#ip-ports\)](#)
- \* [\[InfiniBand\]\(troubleshooting.html#infiniband\)](#)
- \* [\[RDMA over Converged Ethernet \(RoCE\)\]\(troubleshooting.html#rdma-over-converged-ethernet-roce\)](#)

## # Indices and tables

- \* [\[Index\]\(genindex.html\)](#)
- \* [\[Module Index\]\(py-modindex.html\)](#)
- \* [\[Search Page\]\(search.html\)](#)

[\[Next \]\(overview.html "Overview of NCCL"\)](#)

\* \* \*

(C) Copyright 2020, NVIDIA Corporation

Built with [Sphinx](<http://sphinx-doc.org/>) using a

[theme]([https://github.com/rtfd/sphinx\\_rtd\\_theme](https://github.com/rtfd/sphinx_rtd_theme)) provided by [Read the Docs](<https://readthedocs.org>).