

[NCCL](../index.html)

[2.25](https://docs.nvidia.com/deeplearning/sdk/nccl-archived/index.html)

- * [Overview of NCCL](../overview.html)

- * [Setup](../setup.html)

- * [Using NCCL](../usage.html)

- * [Creating a Communicator](communicators.html)

- * [Creating a communicator with options](communicators.html#creating-a-communicator-with-options)

- * [Creating a communicator using multiple ncclUniqueId](communicators.html#creating-a-communicator-using-multiple-nccluniqueids)

- * [Creating more communicators](communicators.html#creating-more-communicators)

- * [Using multiple NCCL communicators concurrently](communicators.html#using-multiple-nccl-communicators-concurrently)

- * [Finalizing a communicator](communicators.html#finalizing-a-communicator)

- * [Destroying a communicator](communicators.html#destroying-a-communicator)

- * [Error handling and communicator abort](communicators.html#error-handling-and-communicator-abort)

- * [Asynchronous errors and error handling](communicators.html#asynchronous-errors-and-error-handling)

- * [Fault Tolerance](communicators.html#fault-tolerance)

- * Collective Operations

- * AllReduce

- * Broadcast

- * Reduce

- * AllGather

- * [ReduceScatter](#)
- * [\[Data Pointers\]\(data.html\)](#)
- * [\[CUDA Stream Semantics\]\(streams.html\)](#)
 - * [\[Mixing Multiple Streams within the same ncclGroupStart/End\(\) group\]\(streams.html#mixing-multiple-streams-within-the-same-ncclgroupstart-end-group\)](#)
- * [\[Group Calls\]\(groups.html\)](#)
 - * [\[Management Of Multiple GPUs From One Thread\]\(groups.html#management-of-multiple-gpus-from-one-thread\)](#)
 - * [\[Aggregated Operations \(2.2 and later\)\]\(groups.html#aggregated-operations-2-2-and-later\)](#)
 - * [\[Nonblocking Group Operation\]\(groups.html#nonblocking-group-operation\)](#)
- * [\[Point-to-point communication\]\(p2p.html\)](#)
 - * [\[Sendrecv\]\(p2p.html#sendrecv\)](#)
 - * [\[One-to-all \(scatter\)\]\(p2p.html#one-to-all-scatter\)](#)
 - * [\[All-to-one \(gather\)\]\(p2p.html#all-to-one-gather\)](#)
 - * [\[All-to-all\]\(p2p.html#all-to-all\)](#)
 - * [\[Neighbor exchange\]\(p2p.html#neighbor-exchange\)](#)
- * [\[Thread Safety\]\(threadsafety.html\)](#)
- * [\[In-place Operations\]\(inplace.html\)](#)
- * [\[Using NCCL with CUDA Graphs\]\(cudagraph.html\)](#)
- * [\[User Buffer Registration\]\(bufferreg.html\)](#)
 - * [\[NVLink Sharp Buffer Registration\]\(bufferreg.html#nvlink-sharp-buffer-registration\)](#)
 - * [\[IB Sharp Buffer Registration\]\(bufferreg.html#ib-sharp-buffer-registration\)](#)
 - * [\[General Buffer Registration\]\(bufferreg.html#general-buffer-registration\)](#)
 - * [\[Memory Allocator\]\(bufferreg.html#memory-allocator\)](#)
- * [\[NCCL API\]\(../api.html\)](#)
 - * [\[Communicator Creation and Management Functions\]\(../api/comms.html\)](#)
 - * [\[ncclGetLastError\]\(../api/comms.html#ncclgetlasterror\)](#)

- * [\[ncclGetErrorString\]\(../api/comms.html#ncclgeterrorstring\)](#)
- * [\[ncclGetVersion\]\(../api/comms.html#ncclgetversion\)](#)
- * [\[ncclGetUniqueId\]\(../api/comms.html#ncclgetuniqueid\)](#)
- * [\[ncclCommInitRank\]\(../api/comms.html#ncclcomminitrank\)](#)
- * [\[ncclCommInitAll\]\(../api/comms.html#ncclcomminitall\)](#)
- * [\[ncclCommInitRankConfig\]\(../api/comms.html#ncclcomminitrankconfig\)](#)
- * [\[ncclCommInitRankScalable\]\(../api/comms.html#ncclcomminitrankscalable\)](#)
- * [\[ncclCommSplit\]\(../api/comms.html#ncclcommsplit\)](#)
- * [\[ncclCommFinalize\]\(../api/comms.html#ncclcommfinalize\)](#)
- * [\[ncclCommDestroy\]\(../api/comms.html#ncclcommdestroy\)](#)
- * [\[ncclCommAbort\]\(../api/comms.html#ncclcommabort\)](#)
- * [\[ncclCommGetAsyncError\]\(../api/comms.html#ncclcommgetasyncerror\)](#)
- * [\[ncclCommCount\]\(../api/comms.html#ncclcommcount\)](#)
- * [\[ncclCommCuDevice\]\(../api/comms.html#ncclcommcudevice\)](#)
- * [\[ncclCommUserRank\]\(../api/comms.html#ncclcommuserrank\)](#)
- * [\[ncclCommRegister\]\(../api/comms.html#ncclcommregister\)](#)
- * [\[ncclCommDeregister\]\(../api/comms.html#ncclcommderegister\)](#)
- * [\[ncclMemAlloc\]\(../api/comms.html#ncclmemalloc\)](#)
- * [\[ncclMemFree\]\(../api/comms.html#ncclmemfree\)](#)
- * [\[Collective Communication Functions\]\(../api/colls.html\)](#)
 - * [\[ncclAllReduce\]\(../api/colls.html#ncclallreduce\)](#)
 - * [\[ncclBroadcast\]\(../api/colls.html#ncclbroadcast\)](#)
 - * [\[ncclReduce\]\(../api/colls.html#ncclreduce\)](#)
 - * [\[ncclAllGather\]\(../api/colls.html#ncclallgather\)](#)
 - * [\[ncclReduceScatter\]\(../api/colls.html#ncclreducescatter\)](#)
- * [\[Group Calls\]\(../api/group.html\)](#)
 - * [\[ncclGroupStart\]\(../api/group.html#ncclgroupstart\)](#)

- * [\[ncclGroupEnd\]\(../api/group.html#ncclgroupend\)](#)
- * [\[ncclGroupSimulateEnd\]\(../api/group.html#ncclgroupsimulateend\)](#)
- * [\[Point To Point Communication Functions\]\(../api/p2p.html\)](#)
- * [\[ncclSend\]\(../api/p2p.html#ncclsend\)](#)
- * [\[ncclRecv\]\(../api/p2p.html#ncclrecv\)](#)
- * [\[Types\]\(../api/types.html\)](#)
- * [\[ncclComm_t\]\(../api/types.html#ncclcomm-t\)](#)
- * [\[ncclResult_t\]\(../api/types.html#ncclresult-t\)](#)
- * [\[ncclDataType_t\]\(../api/types.html#nccldatatype-t\)](#)
- * [\[ncclRedOp_t\]\(../api/types.html#ncclredop-t\)](#)
- * [\[ncclScalarResidence_t\]\(../api/types.html#ncclscalarresidence-t\)](#)
- * [\[ncclConfig_t\]\(../api/types.html#ncclconfig-t\)](#)
- * [\[ncclSimInfo_t\]\(../api/types.html#ncclsiminfo-t\)](#)
- * [\[User Defined Reduction Operators\]\(../api/ops.html\)](#)
- * [\[ncclRedOpCreatePreMulSum\]\(../api/ops.html#ncclredopcreatepremulsum\)](#)
- * [\[ncclRedOpDestroy\]\(../api/ops.html#ncclredopdestroy\)](#)
- * [\[Migrating from NCCL 1 to NCCL 2\]\(../nccl1.html\)](#)
- * [\[Initialization\]\(../nccl1.html#initialization\)](#)
- * [\[Communication\]\(../nccl1.html#communication\)](#)
- * [\[Counts\]\(../nccl1.html#counts\)](#)
- * [\[In-place usage for AllGather and ReduceScatter\]\(../nccl1.html#in-place-usage-for-allgather-and-reducescatter\)](#)
- * [\[AllGather arguments order\]\(../nccl1.html#allgather-arguments-order\)](#)
- * [\[Datatypes\]\(../nccl1.html#datatypes\)](#)
- * [\[Error codes\]\(../nccl1.html#error-codes\)](#)
- * [\[Examples\]\(../examples.html\)](#)
- * [\[Communicator Creation and Destruction\]](#)

Examples](../examples.html#communicator-creation-and-destruction-examples)

- * [Example 1: Single Process, Single Thread, Multiple Devices](../examples.html#example-1-single-process-single-thread-multiple-devices)

- * [Example 2: One Device per Process or Thread](../examples.html#example-2-one-device-per-process-or-thread)

- * [Example 3: Multiple Devices per Thread](../examples.html#example-3-multiple-devices-per-thread)

- * [Example 4: Multiple communicators per device](../examples.html#example-4-multiple-communicators-per-device)

- * [Communication Examples](../examples.html#communication-examples)

- * [Example 1: One Device per Process or Thread](../examples.html#example-1-one-device-per-process-or-thread)

- * [Example 2: Multiple Devices per Thread](../examples.html#example-2-multiple-devices-per-thread)

- * [NCCL and MPI](../mpi.html)

- * [API](../mpi.html#api)

- * [Using multiple devices per process](../mpi.html#using-multiple-devices-per-process)

- * [ReduceScatter operation](../mpi.html#reducescatter-operation)

- * [Send and Receive counts](../mpi.html#send-and-receive-counts)

- * [Other collectives and point-to-point operations](../mpi.html#other-collectives-and-point-to-point-operations)

- * [In-place operations](../mpi.html#in-place-operations)

- * [Using NCCL within an MPI Program](../mpi.html#using-nccl-within-an-mpi-program)

- * [MPI Progress](../mpi.html#mpi-progress)

- * [Inter-GPU Communication with CUDA-aware MPI](../mpi.html#inter-gpu-communication-with-cuda-aware-mpi)

- * [Environment Variables](../env.html)

* [System configuration](../env.html#system-configuration)

* [NCCL_SOCKET_IFNAME](../env.html#nccl-socket-ifname)

* [Values accepted](../env.html#values-accepted)

* [NCCL_SOCKET_FAMILY](../env.html#nccl-socket-family)

* [Values accepted](../env.html#id2)

* [NCCL_SOCKET_RETRY_CNT](../env.html#nccl-socket-retry-cnt)

* [Values accepted](../env.html#id3)

* [NCCL_SOCKET_RETRY_SLEEP_MSEC](../env.html#nccl-socket-retry-sleep-msec)

* [Values accepted](../env.html#id4)

* [NCCL_SOCKET_NTHREADS](../env.html#nccl-socket-nthreads)

* [Values accepted](../env.html#id5)

* [NCCL_NSOCKS_PERTHREAD](../env.html#nccl-nsocks-perthread)

* [Values accepted](../env.html#id6)

* [NCCL_CROSS_NIC](../env.html#nccl-cross-nic)

* [Values accepted](../env.html#id7)

* [NCCL_IB_HCA](../env.html#nccl-ib-hca)

* [Values accepted](../env.html#id8)

* [NCCL_IB_TIMEOUT](../env.html#nccl-ib-timeout)

* [Values accepted](../env.html#id9)

* [NCCL_IB_RETRY_CNT](../env.html#nccl-ib-retry-cnt)

* [Values accepted](../env.html#id10)

* [NCCL_IB_GID_INDEX](../env.html#nccl-ib-gid-index)

* [Values accepted](../env.html#id11)

* [NCCL_IB_ADDR_FAMILY](../env.html#nccl-ib-addr-family)

* [Values accepted](../env.html#id12)

* [NCCL_IB_ADDR_RANGE](../env.html#nccl-ib-addr-range)

* [Values accepted](../env.html#id13)

* [NCCL_IB_ROCE_VERSION_NUM](../env.html#nccl-ib-roce-version-num)

* [Values accepted](../env.html#id14)

* [NCCL_IB_SL](../env.html#nccl-ib-sl)

* [Values accepted](../env.html#id15)

* [NCCL_IB_TC](../env.html#nccl-ib-tc)

* [Values accepted](../env.html#id16)

* [NCCL_IB_FIFO_TC](../env.html#nccl-ib-fifo-tc)

* [Values accepted](../env.html#id17)

* [NCCL_IB_RETURN_ASYNC_EVENTS](../env.html#nccl-ib-return-async-events)

* [Values accepted](../env.html#id18)

* [NCCL_OOB_NET_ENABLE](../env.html#nccl-oob-net-enable)

* [Values accepted](../env.html#id19)

* [NCCL_OOB_NET_IFNAME](../env.html#nccl-oob-net-ifname)

* [Values accepted](../env.html#id20)

* [NCCL_UID_STAGGER_THRESHOLD](../env.html#nccl-uid-stagger-threshold)

* [Values accepted](../env.html#id21)

* [NCCL_UID_STAGGER_RATE](../env.html#nccl-uid-stagger-rate)

* [Values accepted](../env.html#id22)

* [NCCL_NET](../env.html#nccl-net)

* [Values accepted](../env.html#id23)

* [NCCL_NET_PLUGIN](../env.html#nccl-net-plugin)

* [Values accepted](../env.html#id24)

* [NCCL_TUNER_PLUGIN](../env.html#nccl-tuner-plugin)

* [Values accepted](../env.html#id25)

* [NCCL_PROFILER_PLUGIN](../env.html#nccl-profiler-plugin)

* [Values accepted](../env.html#id26)

* [NCCL_IGNORE_CPU_AFFINITY](../env.html#nccl-ignore-cpu-affinity)

- * [Values accepted](../env.html#id27)
- * [NCCL_CONF_FILE](../env.html#nccl-conf-file)
- * [Values accepted](../env.html#id28)
- * [NCCL_DEBUG](../env.html#nccl-debug)
- * [Values accepted](../env.html#id30)
- * [NCCL_DEBUG_FILE](../env.html#nccl-debug-file)
- * [Values accepted](../env.html#id31)
- * [NCCL_DEBUG_SUBSYS](../env.html#nccl-debug-subsys)
- * [Values accepted](../env.html#id32)
- * [NCCL_COLLNET_ENABLE](../env.html#nccl-collnet-enable)
- * [Value accepted](../env.html#value-accepted)
- * [NCCL_COLLNET_NODE_THRESHOLD](../env.html#nccl-collnet-node-threshold)
- * [Value accepted](../env.html#id33)
- * [NCCL_TOPO_FILE](../env.html#nccl-topo-file)
- * [Value accepted](../env.html#id34)
- * [NCCL_TOPO_DUMP_FILE](../env.html#nccl-topo-dump-file)
- * [Value accepted](../env.html#id35)
- * [NCCL_SET_THREAD_NAME](../env.html#nccl-set-thread-name)
- * [Value accepted](../env.html#id36)
- * [Debugging](../env.html#debugging)
- * [NCCL_P2P_DISABLE](../env.html#nccl-p2p-disable)
- * [Values accepted](../env.html#id37)
- * [NCCL_P2P_LEVEL](../env.html#nccl-p2p-level)
- * [Values accepted](../env.html#id38)
- * [Integer Values (Legacy)](../env.html#integer-values-legacy)
- * [NCCL_P2P_DIRECT_DISABLE](../env.html#nccl-p2p-direct-disable)
- * [Values accepted](../env.html#id39)

* [NCCL_SHM_DISABLE](../env.html#nccl-shm-disable)
* [Values accepted](../env.html#id40)

* [NCCL_BUFFSIZE](../env.html#nccl-buffersize)
* [Values accepted](../env.html#id41)

* [NCCL_NTHREADS](../env.html#nccl-nthreads)
* [Values accepted](../env.html#id42)

* [NCCL_MAX_NCHANNELS](../env.html#nccl-max-nchannels)
* [Values accepted](../env.html#id43)

* [NCCL_MIN_NCHANNELS](../env.html#nccl-min-nchannels)
* [Values accepted](../env.html#id44)

* [NCCL_CHECKS_DISABLE](../env.html#nccl-checks-disable)
* [Values accepted](../env.html#id45)

* [NCCL_CHECK_POINTERS](../env.html#nccl-check-pointers)
* [Values accepted](../env.html#id46)

* [NCCL_LAUNCH_MODE](../env.html#nccl-launch-mode)
* [Values accepted](../env.html#id47)

* [NCCL_IB_DISABLE](../env.html#nccl-ib-disable)
* [Values accepted](../env.html#id48)

* [NCCL_IB_AR_THRESHOLD](../env.html#nccl-ib-ar-threshold)
* [Values accepted](../env.html#id49)

* [NCCL_IB_QPS_PER_CONNECTION](../env.html#nccl-ib-qps-per-connection)
* [Values accepted](../env.html#id50)

* [NCCL_IB_SPLIT_DATA_ON_QPS](../env.html#nccl-ib-split-data-on-qps)
* [Values accepted](../env.html#id51)

* [NCCL_IB_CUDA_SUPPORT](../env.html#nccl-ib-cuda-support)
* [Values accepted](../env.html#id52)

* [NCCL_IB_PCI_RELAXED_ORDERING](../env.html#nccl-ib-pci-relaxed-ordering)

* [Values accepted](../env.html#id53)

* [NCCL_IB_ADAPTIVE_ROUTING](../env.html#nccl-ib-adaptive-routing)

* [Values accepted](../env.html#id54)

* [NCCL_IB_ECE_ENABLE](../env.html#nccl-ib-ece-enable)

* [Values accepted](../env.html#id55)

* [NCCL_MEM_SYNC_DOMAIN](../env.html#nccl-mem-sync-domain)

* [Values accepted](../env.html#id56)

* [NCCL_CUMEM_ENABLE](../env.html#nccl-cumem-enable)

* [Values accepted](../env.html#id57)

* [NCCL_CUMEM_HOST_ENABLE](../env.html#nccl-cumem-host-enable)

* [Values accepted](../env.html#id58)

* [NCCL_NET_GDR_LEVEL (formerly

NCCL_IB_GDR_LEVEL)](../env.html#nccl-net-gdr-level-formerly-nccl-ib-gdr-level)

* [Values accepted](../env.html#id59)

* [Integer Values (Legacy)](../env.html#id60)

* [NCCL_NET_GDR_READ](../env.html#nccl-net-gdr-read)

* [Values accepted](../env.html#id61)

* [NCCL_NET_SHARED_BUFFERS](../env.html#nccl-net-shared-buffers)

* [Value accepted](../env.html#id62)

* [NCCL_NET_SHARED_COMMS](../env.html#nccl-net-shared-comms)

* [Value accepted](../env.html#id63)

* [NCCL_SINGLE_RING_THRESHOLD](../env.html#nccl-single-ring-threshold)

* [Values accepted](../env.html#id64)

* [NCCL_LL_THRESHOLD](../env.html#nccl-ll-threshold)

* [Values accepted](../env.html#id65)

* [NCCL_TREE_THRESHOLD](../env.html#nccl-tree-threshold)

* [Values accepted](../env.html#id66)

* [NCCL_ALGO](../env.html#nccl-algo)

* [Values accepted](../env.html#id67)

* [NCCL_PROTO](../env.html#nccl-proto)

* [Values accepted](../env.html#id68)

* [NCCL_NVX_DISABLE](../env.html#nccl-nvx-disable)

* [Value accepted](../env.html#id69)

* [NCCL_P2P_DISABLE](../env.html#nccl-p2p-disable)

* [Value accepted](../env.html#id70)

* [NCCL_P2P_P2P_LEVEL](../env.html#nccl-p2p-p2p-level)

* [Value accepted](../env.html#id71)

* [NCCL_RUNTIME_CONNECT](../env.html#nccl-runtime-connect)

* [Value accepted](../env.html#id72)

* [NCCL_GRAPH_REGISTER](../env.html#nccl-graph-register)

* [Value accepted](../env.html#id74)

* [NCCL_LOCAL_REGISTER](../env.html#nccl-local-register)

* [Value accepted](../env.html#id75)

* [NCCL_LEGACY_CUDA_REGISTER](../env.html#nccl-legacy-cuda-register)

* [Value accepted](../env.html#id76)

* [NCCL_SET_STACK_SIZE](../env.html#nccl-set-stack-size)

* [Value accepted](../env.html#id77)

* [NCCL_GRAPH_MIXING_SUPPORT](../env.html#nccl-graph-mixing-support)

* [Value accepted](../env.html#id79)

* [NCCL_DMABUF_ENABLE](../env.html#nccl-dmabuf-enable)

* [Value accepted](../env.html#id80)

* [NCCL_P2P_NET_CHUNKSIZE](../env.html#nccl-p2p-net-chunksize)

* [Values accepted](../env.html#id81)

* [NCCL_P2P_LL_THRESHOLD](../env.html#nccl-p2p-ll-threshold)

- * [\[Values accepted\]\(../env.html#id82\)](#)
- * [\[NCCL_ALLOC_P2P_NET_LL_BUFFERS\]\(../env.html#nccl-alloc-p2p-net-ll-buffers\)](#)
- * [\[Values accepted\]\(../env.html#id83\)](#)
- * [\[NCCL_COMM_BLOCKING\]\(../env.html#nccl-comm-blocking\)](#)
- * [\[Values accepted\]\(../env.html#id84\)](#)
- * [\[NCCL_CGA_CLUSTER_SIZE\]\(../env.html#nccl-cga-cluster-size\)](#)
- * [\[Values accepted\]\(../env.html#id85\)](#)
- * [\[NCCL_MAX_CTAS\]\(../env.html#nccl-max-ctas\)](#)
- * [\[Values accepted\]\(../env.html#id86\)](#)
- * [\[NCCL_MIN_CTAS\]\(../env.html#nccl-min-ctas\)](#)
- * [\[Values accepted\]\(../env.html#id87\)](#)
- * [\[NCCL_NVLS_ENABLE\]\(../env.html#nccl-nvls-enable\)](#)
- * [\[Values accepted\]\(../env.html#id88\)](#)
- * [\[NCCL_IB_MERGE_NICS\]\(../env.html#nccl-ib-merge-nics\)](#)
- * [\[Values accepted\]\(../env.html#id89\)](#)
- * [\[NCCL_MNNVL_ENABLE\]\(../env.html#nccl-mnnvl-enable\)](#)
- * [\[Values accepted\]\(../env.html#id90\)](#)
- * [\[NCCL_RAS_ENABLE\]\(../env.html#nccl-ras-enable\)](#)
- * [\[Values accepted\]\(../env.html#id91\)](#)
- * [\[NCCL_RAS_ADDR\]\(../env.html#nccl-ras-addr\)](#)
- * [\[Values accepted\]\(../env.html#id92\)](#)
- * [\[NCCL_RAS_TIMEOUT_FACTOR\]\(../env.html#nccl-ras-timeout-factor\)](#)
- * [\[Values accepted\]\(../env.html#id93\)](#)
- * [\[Troubleshooting\]\(../troubleshooting.html\)](#)
- * [\[Errors\]\(../troubleshooting.html#errors\)](#)
- * [\[RAS\]\(../troubleshooting.html#ras\)](#)
- * [\[RAS\]\(../troubleshooting/ras.html\)](#)

- * [\[Principle of Operation\]\(../troubleshooting/ras.html#principle-of-operation\)](#)
- * [\[RAS Queries\]\(../troubleshooting/ras.html#ras-queries\)](#)
- * [\[Sample Output\]\(../troubleshooting/ras.html#sample-output\)](#)
- * [\[GPU Direct\]\(../troubleshooting.html#gpu-direct\)](#)
- * [\[GPU-to-GPU communication\]\(../troubleshooting.html#gpu-to-gpu-communication\)](#)
- * [\[GPU-to-NIC communication\]\(../troubleshooting.html#gpu-to-nic-communication\)](#)
- * [\[PCI Access Control Services \(ACS\)\]\(../troubleshooting.html#pci-access-control-services-ac\)](#)
- * [\[Topology detection\]\(../troubleshooting.html#topology-detection\)](#)
- * [\[Shared memory\]\(../troubleshooting.html#shared-memory\)](#)
- * [\[Docker\]\(../troubleshooting.html#docker\)](#)
- * [\[Systemd\]\(../troubleshooting.html#systemd\)](#)
- * [\[Networking issues\]\(../troubleshooting.html#networking-issues\)](#)
- * [\[IP Network Interfaces\]\(../troubleshooting.html#ip-network-interfaces\)](#)
- * [\[IP Ports\]\(../troubleshooting.html#ip-ports\)](#)
- * [\[InfiniBand\]\(../troubleshooting.html#infiniband\)](#)
- * [\[RDMA over Converged Ethernet \(RoCE\)\]\(../troubleshooting.html#rdma-over-converged-ethernet-roce\)](#)

[__\[NCCL\]\(../index.html\)](#)

- * [\[Docs\]\(../index.html\)](#) »
- * [\[Using NCCL\]\(../usage.html\)](#) »
- * [Collective Operations](#)
- * [\[View page source\]\(../_sources/usage/collectives.rst.txt\)](#)

* * *

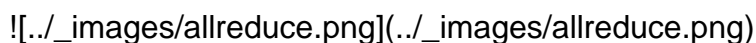
Collective Operations¶

Collective operations have to be called for each rank (hence CUDA device), using the same count and the same datatype, to form a complete collective operation. Failure to do so will result in undefined behavior, including hangs, crashes, or data corruption.

AllReduce¶

The AllReduce operation performs reductions on data (for example, sum, min, max) across devices and stores the result in the receive buffer of every rank.

In a `_sum_` allreduce operation between `_k_` ranks, each rank will provide an array in of N values, and receive identical results in array out of N values, where $out[i] = in_0[i] + in_1[i] + \dots + in_{(k-1)}[i]$.

The diagram shows a sequence of ranks (0 to k-1) each providing an input array of N values. These are combined via an allreduce operation to produce a single output array of N values that is identical on all ranks.

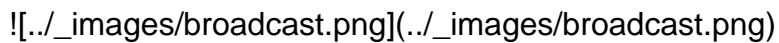
All-Reduce operation: each rank receives the reduction of input values across ranks.

Related links: [`ncclAllReduce()`](../api/colls.html#c.ncclAllReduce "ncclAllReduce").

Broadcast¶

The Broadcast operation copies an N-element buffer from the root rank to all

the ranks.

![_images/broadcast.png](../_images/broadcast.png)

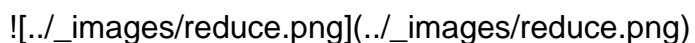
Broadcast operation: all ranks receive data from a "root" rank.

Important note: The root argument is one of the ranks, not a device number, and is therefore impacted by a different rank to device mapping.

Related links: [ncclBroadcast()](../api/colls.html#c.ncclBroadcast "ncclBroadcast").

Reduce¶

The Reduce operation performs the same operation as AllReduce, but stores the result only in the receive buffer of a specified root rank.

![_images/reduce.png](../_images/reduce.png)

Reduce operation: one rank receives the reduction of input values across ranks.

Important note: The root argument is one of the ranks (not a device number), and is therefore impacted by a different rank to device mapping.

Note: A Reduce, followed by a Broadcast, is equivalent to the AllReduce operation.

Related links: [`ncclReduce()`](../api/colls.html#c.ncclReduce "ncclReduce").

AllGather¶

The AllGather operation gathers N values from k ranks into an output buffer of size $k*N$, and distributes that result to all ranks.

The output is ordered by the rank index. The AllGather operation is therefore impacted by a different rank to device mapping.

![_images/allgather.png](../_images/allgather.png)

AllGather operation: each rank receives the aggregation of data from all ranks in the order of the ranks.

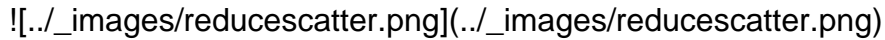
Note: Executing ReduceScatter, followed by AllGather, is equivalent to the AllReduce operation.

Related links: [`ncclAllGather()`](../api/colls.html#c.ncclAllGather "ncclAllGather").

ReduceScatter¶

The ReduceScatter operation performs the same operation as Reduce, except that the result is scattered in equal-sized blocks between ranks, each rank getting a chunk of data based on its rank index.

The ReduceScatter operation is impacted by a different rank to device mapping since the ranks determine the data layout.



Reduce-Scatter operation: input values are reduced across ranks, with each rank receiving a subpart of the result.

Related links: `[`ncclReduceScatter()`](../api/colls.html#c.ncclReduceScatter "ncclReduceScatter")`

[\[Next \]\(data.html "Data Pointers"\)](#) [\[Previous\]\(communicators.html "Creating a Communicator"\)](#)

* * *

(C) Copyright 2020, NVIDIA Corporation

Built with [\[Sphinx\]](http://sphinx-doc.org/) using a [\[theme\]](https://github.com/rtfd/sphinx_rtd_theme) provided by [\[Read the Docs\]](https://readthedocs.org).