

[ NCCL ](../index.html)

[2.25](https://docs.nvidia.com/deeplearning/sdk/nccl-archived/index.html)

- \* [Overview of NCCL](../overview.html)

- \* [Setup](../setup.html)

- \* [Using NCCL](../usage.html)

- \* [Creating a Communicator](communicators.html)

- \* [Creating a communicator with options](communicators.html#creating-a-communicator-with-options)

- \* [Creating a communicator using multiple ncclUniqueId](communicators.html#creating-a-communicator-using-multiple-nccluniqueids)

- \* [Creating more communicators](communicators.html#creating-more-communicators)

- \* [Using multiple NCCL communicators concurrently](communicators.html#using-multiple-nccl-communicators-concurrently)

- \* [Finalizing a communicator](communicators.html#finalizing-a-communicator)

- \* [Destroying a communicator](communicators.html#destroying-a-communicator)

- \* [Error handling and communicator abort](communicators.html#error-handling-and-communicator-abort)

- \* [Asynchronous errors and error handling](communicators.html#asynchronous-errors-and-error-handling)

- \* [Fault Tolerance](communicators.html#fault-tolerance)

- \* [Collective Operations](collectives.html)

- \* [AllReduce](collectives.html#allreduce)

- \* [Broadcast](collectives.html#broadcast)

- \* [Reduce](collectives.html#reduce)

- \* [AllGather](collectives.html#allgather)

- \* [\[ReduceScatter\]\(collectives.html#reducescatter\)](#)
- \* [\[Data Pointers\]\(data.html\)](#)
- \* [\[CUDA Stream Semantics\]\(streams.html\)](#)
  - \* [\[Mixing Multiple Streams within the same ncclGroupStart/End\(\) group\]\(streams.html#mixing-multiple-streams-within-the-same-ncclgroupstart-end-group\)](#)
- \* Group Calls
  - \* Management Of Multiple GPUs From One Thread
  - \* Aggregated Operations (2.2 and later)
  - \* Nonblocking Group Operation
- \* [\[Point-to-point communication\]\(p2p.html\)](#)
  - \* [\[Sendrecv\]\(p2p.html#sendrecv\)](#)
  - \* [\[One-to-all \(scatter\)\]\(p2p.html#one-to-all-scatter\)](#)
  - \* [\[All-to-one \(gather\)\]\(p2p.html#all-to-one-gather\)](#)
  - \* [\[All-to-all\]\(p2p.html#all-to-all\)](#)
  - \* [\[Neighbor exchange\]\(p2p.html#neighbor-exchange\)](#)
- \* [\[Thread Safety\]\(threadsafety.html\)](#)
- \* [\[In-place Operations\]\(inplace.html\)](#)
- \* [\[Using NCCL with CUDA Graphs\]\(cudagraph.html\)](#)
- \* [\[User Buffer Registration\]\(bufferreg.html\)](#)
  - \* [\[NVLink Sharp Buffer Registration\]\(bufferreg.html#nvlink-sharp-buffer-registration\)](#)
  - \* [\[IB Sharp Buffer Registration\]\(bufferreg.html#ib-sharp-buffer-registration\)](#)
  - \* [\[General Buffer Registration\]\(bufferreg.html#general-buffer-registration\)](#)
  - \* [\[Memory Allocator\]\(bufferreg.html#memory-allocator\)](#)
- \* [\[NCCL API\]\(../api.html\)](#)
  - \* [\[Communicator Creation and Management Functions\]\(../api/comms.html\)](#)
  - \* [\[ncclGetLastError\]\(../api/comms.html#ncclgetlasterror\)](#)
  - \* [\[ncclGetErrorString\]\(../api/comms.html#ncclgeterrorstring\)](#)

- \* [\[ncclGetVersion\]\(../api/comms.html#ncclgetversion\)](#)
- \* [\[ncclGetUniqueId\]\(../api/comms.html#ncclgetuniqueid\)](#)
- \* [\[ncclCommInitRank\]\(../api/comms.html#ncclcomminitrank\)](#)
- \* [\[ncclCommInitAll\]\(../api/comms.html#ncclcomminitall\)](#)
- \* [\[ncclCommInitRankConfig\]\(../api/comms.html#ncclcomminitrankconfig\)](#)
- \* [\[ncclCommInitRankScalable\]\(../api/comms.html#ncclcomminitrankscalable\)](#)
- \* [\[ncclCommSplit\]\(../api/comms.html#ncclcommsplit\)](#)
- \* [\[ncclCommFinalize\]\(../api/comms.html#ncclcommfinalize\)](#)
- \* [\[ncclCommDestroy\]\(../api/comms.html#ncclcommdestroy\)](#)
- \* [\[ncclCommAbort\]\(../api/comms.html#ncclcommabort\)](#)
- \* [\[ncclCommGetAsyncError\]\(../api/comms.html#ncclcommgetasyncerror\)](#)
- \* [\[ncclCommCount\]\(../api/comms.html#ncclcommcount\)](#)
- \* [\[ncclCommCuDevice\]\(../api/comms.html#ncclcommcudevice\)](#)
- \* [\[ncclCommUserRank\]\(../api/comms.html#ncclcommuserrank\)](#)
- \* [\[ncclCommRegister\]\(../api/comms.html#ncclcommregister\)](#)
- \* [\[ncclCommDeregister\]\(../api/comms.html#ncclcommderegister\)](#)
- \* [\[ncclMemAlloc\]\(../api/comms.html#ncclmemalloc\)](#)
- \* [\[ncclMemFree\]\(../api/comms.html#ncclmemfree\)](#)
- \* [\[Collective Communication Functions\]\(../api/colls.html\)](#)
  - \* [\[ncclAllReduce\]\(../api/colls.html#ncclallreduce\)](#)
  - \* [\[ncclBroadcast\]\(../api/colls.html#ncclbroadcast\)](#)
  - \* [\[ncclReduce\]\(../api/colls.html#ncclreduce\)](#)
  - \* [\[ncclAllGather\]\(../api/colls.html#ncclallgather\)](#)
  - \* [\[ncclReduceScatter\]\(../api/colls.html#ncclreducescatter\)](#)
- \* [\[Group Calls\]\(../api/group.html\)](#)
  - \* [\[ncclGroupStart\]\(../api/group.html#ncclgroupstart\)](#)
  - \* [\[ncclGroupEnd\]\(../api/group.html#ncclgroupend\)](#)

- \* [\[ncclGroupSimulateEnd\]\(../api/group.html#ncclgroupsimulateend\)](#)
- \* [\[Point To Point Communication Functions\]\(../api/p2p.html\)](#)
- \* [\[ncclSend\]\(../api/p2p.html#ncclsend\)](#)
- \* [\[ncclRecv\]\(../api/p2p.html#ncclrecv\)](#)
- \* [\[Types\]\(../api/types.html\)](#)
- \* [\[ncclComm\\_t\]\(../api/types.html#ncclcomm-t\)](#)
- \* [\[ncclResult\\_t\]\(../api/types.html#ncclresult-t\)](#)
- \* [\[ncclDataType\\_t\]\(../api/types.html#nccldatatype-t\)](#)
- \* [\[ncclRedOp\\_t\]\(../api/types.html#ncclredop-t\)](#)
- \* [\[ncclScalarResidence\\_t\]\(../api/types.html#ncclscalarresidence-t\)](#)
- \* [\[ncclConfig\\_t\]\(../api/types.html#ncclconfig-t\)](#)
- \* [\[ncclSimInfo\\_t\]\(../api/types.html#ncclsiminfo-t\)](#)
- \* [\[User Defined Reduction Operators\]\(../api/ops.html\)](#)
- \* [\[ncclRedOpCreatePreMulSum\]\(../api/ops.html#ncclredopcreatepremulsum\)](#)
- \* [\[ncclRedOpDestroy\]\(../api/ops.html#ncclredopdestroy\)](#)
- \* [\[Migrating from NCCL 1 to NCCL 2\]\(../nccl1.html\)](#)
- \* [\[Initialization\]\(../nccl1.html#initialization\)](#)
- \* [\[Communication\]\(../nccl1.html#communication\)](#)
- \* [\[Counts\]\(../nccl1.html#counts\)](#)
- \* [\[In-place usage for AllGather and ReduceScatter\]\(../nccl1.html#in-place-usage-for-allgather-and-reducescatter\)](#)
- \* [\[AllGather arguments order\]\(../nccl1.html#allgather-arguments-order\)](#)
- \* [\[Datatypes\]\(../nccl1.html#datatypes\)](#)
- \* [\[Error codes\]\(../nccl1.html#error-codes\)](#)
- \* [\[Examples\]\(../examples.html\)](#)
- \* [\[Communicator Creation and Destruction Examples\]\(../examples.html#communicator-creation-and-destruction-examples\)](#)

- \* [Example 1: Single Process, Single Thread, Multiple Devices](../examples.html#example-1-single-process-single-thread-multiple-devices)
- \* [Example 2: One Device per Process or Thread](../examples.html#example-2-one-device-per-process-or-thread)
- \* [Example 3: Multiple Devices per Thread](../examples.html#example-3-multiple-devices-per-thread)
- \* [Example 4: Multiple communicators per device](../examples.html#example-4-multiple-communicators-per-device)
- \* [Communication Examples](../examples.html#communication-examples)
  - \* [Example 1: One Device per Process or Thread](../examples.html#example-1-one-device-per-process-or-thread)
  - \* [Example 2: Multiple Devices per Thread](../examples.html#example-2-multiple-devices-per-thread)
- \* [NCCL and MPI](../mpi.html)
  - \* [API](../mpi.html#api)
    - \* [Using multiple devices per process](../mpi.html#using-multiple-devices-per-process)
    - \* [ReduceScatter operation](../mpi.html#reducescatter-operation)
    - \* [Send and Receive counts](../mpi.html#send-and-receive-counts)
    - \* [Other collectives and point-to-point operations](../mpi.html#other-collectives-and-point-to-point-operations)
    - \* [In-place operations](../mpi.html#in-place-operations)
  - \* [Using NCCL within an MPI Program](../mpi.html#using-nccl-within-an-mpi-program)
  - \* [MPI Progress](../mpi.html#mpi-progress)
  - \* [Inter-GPU Communication with CUDA-aware MPI](../mpi.html#inter-gpu-communication-with-cuda-aware-mpi)
- \* [Environment Variables](../env.html)
  - \* [System configuration](../env.html#system-configuration)

\* [NCCL\_SOCKET\_IFNAME](../env.html#nccl-socket-ifname)  
\* [Values accepted](../env.html#values-accepted)

\* [NCCL\_SOCKET\_FAMILY](../env.html#nccl-socket-family)  
\* [Values accepted](../env.html#id2)

\* [NCCL\_SOCKET\_RETRY\_CNT](../env.html#nccl-socket-retry-cnt)  
\* [Values accepted](../env.html#id3)

\* [NCCL\_SOCKET\_RETRY\_SLEEP\_MSEC](../env.html#nccl-socket-retry-sleep-msec)  
\* [Values accepted](../env.html#id4)

\* [NCCL\_SOCKET\_NTHREADS](../env.html#nccl-socket-nthreads)  
\* [Values accepted](../env.html#id5)

\* [NCCL\_NSOCKS\_PERTHREAD](../env.html#nccl-nsocks-perthread)  
\* [Values accepted](../env.html#id6)

\* [NCCL\_CROSS\_NIC](../env.html#nccl-cross-nic)  
\* [Values accepted](../env.html#id7)

\* [NCCL\_IB\_HCA](../env.html#nccl-ib-hca)  
\* [Values accepted](../env.html#id8)

\* [NCCL\_IB\_TIMEOUT](../env.html#nccl-ib-timeout)  
\* [Values accepted](../env.html#id9)

\* [NCCL\_IB\_RETRY\_CNT](../env.html#nccl-ib-retry-cnt)  
\* [Values accepted](../env.html#id10)

\* [NCCL\_IB\_GID\_INDEX](../env.html#nccl-ib-gid-index)  
\* [Values accepted](../env.html#id11)

\* [NCCL\_IB\_ADDR\_FAMILY](../env.html#nccl-ib-addr-family)  
\* [Values accepted](../env.html#id12)

\* [NCCL\_IB\_ADDR\_RANGE](../env.html#nccl-ib-addr-range)  
\* [Values accepted](../env.html#id13)

\* [NCCL\_IB\_ROCE\_VERSION\_NUM](../env.html#nccl-ib-roce-version-num)

\* [Values accepted](../env.html#id14)

\* [NCCL\_IB\_SL](../env.html#nccl-ib-sl)

\* [Values accepted](../env.html#id15)

\* [NCCL\_IB\_TC](../env.html#nccl-ib-tc)

\* [Values accepted](../env.html#id16)

\* [NCCL\_IB\_FIFO\_TC](../env.html#nccl-ib-fifo-tc)

\* [Values accepted](../env.html#id17)

\* [NCCL\_IB\_RETURN\_ASYNC\_EVENTS](../env.html#nccl-ib-return-async-events)

\* [Values accepted](../env.html#id18)

\* [NCCL\_OOB\_NET\_ENABLE](../env.html#nccl-oob-net-enable)

\* [Values accepted](../env.html#id19)

\* [NCCL\_OOB\_NET\_IFNAME](../env.html#nccl-oob-net-ifname)

\* [Values accepted](../env.html#id20)

\* [NCCL\_UID\_STAGGER\_THRESHOLD](../env.html#nccl-uid-stagger-threshold)

\* [Values accepted](../env.html#id21)

\* [NCCL\_UID\_STAGGER\_RATE](../env.html#nccl-uid-stagger-rate)

\* [Values accepted](../env.html#id22)

\* [NCCL\_NET](../env.html#nccl-net)

\* [Values accepted](../env.html#id23)

\* [NCCL\_NET\_PLUGIN](../env.html#nccl-net-plugin)

\* [Values accepted](../env.html#id24)

\* [NCCL\_TUNER\_PLUGIN](../env.html#nccl-tuner-plugin)

\* [Values accepted](../env.html#id25)

\* [NCCL\_PROFILER\_PLUGIN](../env.html#nccl-profiler-plugin)

\* [Values accepted](../env.html#id26)

\* [NCCL\_IGNORE\_CPU\_AFFINITY](../env.html#nccl-ignore-cpu-affinity)

\* [Values accepted](../env.html#id27)

- \* [NCCL\_CONF\_FILE](../env.html#nccl-conf-file)
  - \* [Values accepted](../env.html#id28)
- \* [NCCL\_DEBUG](../env.html#nccl-debug)
  - \* [Values accepted](../env.html#id30)
- \* [NCCL\_DEBUG\_FILE](../env.html#nccl-debug-file)
  - \* [Values accepted](../env.html#id31)
- \* [NCCL\_DEBUG\_SUBSYS](../env.html#nccl-debug-subsys)
  - \* [Values accepted](../env.html#id32)
- \* [NCCL\_COLLNET\_ENABLE](../env.html#nccl-collnet-enable)
  - \* [Value accepted](../env.html#value-accepted)
- \* [NCCL\_COLLNET\_NODE\_THRESHOLD](../env.html#nccl-collnet-node-threshold)
  - \* [Value accepted](../env.html#id33)
- \* [NCCL\_TOPO\_FILE](../env.html#nccl-topo-file)
  - \* [Value accepted](../env.html#id34)
- \* [NCCL\_TOPO\_DUMP\_FILE](../env.html#nccl-topo-dump-file)
  - \* [Value accepted](../env.html#id35)
- \* [NCCL\_SET\_THREAD\_NAME](../env.html#nccl-set-thread-name)
  - \* [Value accepted](../env.html#id36)
- \* [Debugging](../env.html#debugging)
- \* [NCCL\_P2P\_DISABLE](../env.html#nccl-p2p-disable)
  - \* [Values accepted](../env.html#id37)
- \* [NCCL\_P2P\_LEVEL](../env.html#nccl-p2p-level)
  - \* [Values accepted](../env.html#id38)
  - \* [Integer Values (Legacy)](../env.html#integer-values-legacy)
- \* [NCCL\_P2P\_DIRECT\_DISABLE](../env.html#nccl-p2p-direct-disable)
  - \* [Values accepted](../env.html#id39)
- \* [NCCL\_SHM\_DISABLE](../env.html#nccl-shm-disable)



\* [Values accepted](../env.html#id40)

\* [NCCL\_BUFFSIZE](../env.html#nccl-buffersize)

\* [Values accepted](../env.html#id41)

\* [NCCL\_NTHREADS](../env.html#nccl-nthreads)

\* [Values accepted](../env.html#id42)

\* [NCCL\_MAX\_NCHANNELS](../env.html#nccl-max-nchannels)

\* [Values accepted](../env.html#id43)

\* [NCCL\_MIN\_NCHANNELS](../env.html#nccl-min-nchannels)

\* [Values accepted](../env.html#id44)

\* [NCCL\_CHECKS\_DISABLE](../env.html#nccl-checks-disable)

\* [Values accepted](../env.html#id45)

\* [NCCL\_CHECK\_POINTERS](../env.html#nccl-check-pointers)

\* [Values accepted](../env.html#id46)

\* [NCCL\_LAUNCH\_MODE](../env.html#nccl-launch-mode)

\* [Values accepted](../env.html#id47)

\* [NCCL\_IB\_DISABLE](../env.html#nccl-ib-disable)

\* [Values accepted](../env.html#id48)

\* [NCCL\_IB\_AR\_THRESHOLD](../env.html#nccl-ib-ar-threshold)

\* [Values accepted](../env.html#id49)

\* [NCCL\_IB\_QPS\_PER\_CONNECTION](../env.html#nccl-ib-qps-per-connection)

\* [Values accepted](../env.html#id50)

\* [NCCL\_IB\_SPLIT\_DATA\_ON\_QPS](../env.html#nccl-ib-split-data-on-qps)

\* [Values accepted](../env.html#id51)

\* [NCCL\_IB\_CUDA\_SUPPORT](../env.html#nccl-ib-cuda-support)

\* [Values accepted](../env.html#id52)

\* [NCCL\_IB\_PCI\_RELAXED\_ORDERING](../env.html#nccl-ib-pci-relaxed-ordering)

\* [Values accepted](../env.html#id53)

\* [NCCL\_IB\_ADAPTIVE\_ROUTING](../env.html#nccl-ib-adaptive-routing)

\* [Values accepted](../env.html#id54)

\* [NCCL\_IB\_ECE\_ENABLE](../env.html#nccl-ib-ece-enable)

\* [Values accepted](../env.html#id55)

\* [NCCL\_MEM\_SYNC\_DOMAIN](../env.html#nccl-mem-sync-domain)

\* [Values accepted](../env.html#id56)

\* [NCCL\_CUMEM\_ENABLE](../env.html#nccl-cumem-enable)

\* [Values accepted](../env.html#id57)

\* [NCCL\_CUMEM\_HOST\_ENABLE](../env.html#nccl-cumem-host-enable)

\* [Values accepted](../env.html#id58)

\* [NCCL\_NET\_GDR\_LEVEL (formerly

NCCL\_IB\_GDR\_LEVEL)](../env.html#nccl-net-gdr-level-formerly-nccl-ib-gdr-level)

\* [Values accepted](../env.html#id59)

\* [Integer Values (Legacy)](../env.html#id60)

\* [NCCL\_NET\_GDR\_READ](../env.html#nccl-net-gdr-read)

\* [Values accepted](../env.html#id61)

\* [NCCL\_NET\_SHARED\_BUFFERS](../env.html#nccl-net-shared-buffers)

\* [Value accepted](../env.html#id62)

\* [NCCL\_NET\_SHARED\_COMMS](../env.html#nccl-net-shared-comms)

\* [Value accepted](../env.html#id63)

\* [NCCL\_SINGLE\_RING\_THRESHOLD](../env.html#nccl-single-ring-threshold)

\* [Values accepted](../env.html#id64)

\* [NCCL\_LL\_THRESHOLD](../env.html#nccl-ll-threshold)

\* [Values accepted](../env.html#id65)

\* [NCCL\_TREE\_THRESHOLD](../env.html#nccl-tree-threshold)

\* [Values accepted](../env.html#id66)

\* [NCCL\_ALGO](../env.html#nccl-algo)

\* [Values accepted](../env.html#id67)

\* [NCCL\_PROTO](../env.html#nccl-proto)

\* [Values accepted](../env.html#id68)

\* [NCCL\_NV\_B\_DISABLE](../env.html#nccl-nvb-disable)

\* [Value accepted](../env.html#id69)

\* [NCCL\_PXN\_DISABLE](../env.html#nccl-pxn-disable)

\* [Value accepted](../env.html#id70)

\* [NCCL\_P2P\_PXN\_LEVEL](../env.html#nccl-p2p-pxn-level)

\* [Value accepted](../env.html#id71)

\* [NCCL\_RUNTIME\_CONNECT](../env.html#nccl-runtime-connect)

\* [Value accepted](../env.html#id72)

\* [NCCL\_GRAPH\_REGISTER](../env.html#nccl-graph-register)

\* [Value accepted](../env.html#id74)

\* [NCCL\_LOCAL\_REGISTER](../env.html#nccl-local-register)

\* [Value accepted](../env.html#id75)

\* [NCCL\_LEGACY\_CUDA\_REGISTER](../env.html#nccl-legacy-cuda-register)

\* [Value accepted](../env.html#id76)

\* [NCCL\_SET\_STACK\_SIZE](../env.html#nccl-set-stack-size)

\* [Value accepted](../env.html#id77)

\* [NCCL\_GRAPH\_MIXING\_SUPPORT](../env.html#nccl-graph-mixing-support)

\* [Value accepted](../env.html#id79)

\* [NCCL\_DMABUF\_ENABLE](../env.html#nccl-dmabuf-enable)

\* [Value accepted](../env.html#id80)

\* [NCCL\_P2P\_NET\_CHUNKSIZE](../env.html#nccl-p2p-net-chunksize)

\* [Values accepted](../env.html#id81)

\* [NCCL\_P2P\_LL\_THRESHOLD](../env.html#nccl-p2p-ll-threshold)

\* [Values accepted](../env.html#id82)

- \* [\[NCCL\\_ALLOC\\_P2P\\_NET\\_LL\\_BUFFERS\]\(../env.html#nccl-alloc-p2p-net-ll-buffers\)](#)
  - \* [\[Values accepted\]\(../env.html#id83\)](#)
- \* [\[NCCL\\_COMM\\_BLOCKING\]\(../env.html#nccl-comm-blocking\)](#)
  - \* [\[Values accepted\]\(../env.html#id84\)](#)
- \* [\[NCCL\\_CGA\\_CLUSTER\\_SIZE\]\(../env.html#nccl-cga-cluster-size\)](#)
  - \* [\[Values accepted\]\(../env.html#id85\)](#)
- \* [\[NCCL\\_MAX\\_CTAS\]\(../env.html#nccl-max-ctas\)](#)
  - \* [\[Values accepted\]\(../env.html#id86\)](#)
- \* [\[NCCL\\_MIN\\_CTAS\]\(../env.html#nccl-min-ctas\)](#)
  - \* [\[Values accepted\]\(../env.html#id87\)](#)
- \* [\[NCCL\\_NVLS\\_ENABLE\]\(../env.html#nccl-nvls-enable\)](#)
  - \* [\[Values accepted\]\(../env.html#id88\)](#)
- \* [\[NCCL\\_IB\\_MERGE\\_NICS\]\(../env.html#nccl-ib-merge-nics\)](#)
  - \* [\[Values accepted\]\(../env.html#id89\)](#)
- \* [\[NCCL\\_MNNVL\\_ENABLE\]\(../env.html#nccl-mnnvl-enable\)](#)
  - \* [\[Values accepted\]\(../env.html#id90\)](#)
- \* [\[NCCL\\_RAS\\_ENABLE\]\(../env.html#nccl-ras-enable\)](#)
  - \* [\[Values accepted\]\(../env.html#id91\)](#)
- \* [\[NCCL\\_RAS\\_ADDR\]\(../env.html#nccl-ras-addr\)](#)
  - \* [\[Values accepted\]\(../env.html#id92\)](#)
- \* [\[NCCL\\_RAS\\_TIMEOUT\\_FACTOR\]\(../env.html#nccl-ras-timeout-factor\)](#)
  - \* [\[Values accepted\]\(../env.html#id93\)](#)
- \* [\[Troubleshooting\]\(../troubleshooting.html\)](#)
  - \* [\[Errors\]\(../troubleshooting.html#errors\)](#)
  - \* [\[RAS\]\(../troubleshooting.html#ras\)](#)
    - \* [\[RAS\]\(../troubleshooting/ras.html\)](#)
      - \* [\[Principle of Operation\]\(../troubleshooting/ras.html#principle-of-operation\)](#)

- \* [\[RAS Queries\]\(../troubleshooting/ras.html#ras-queries\)](#)
- \* [\[Sample Output\]\(../troubleshooting/ras.html#sample-output\)](#)
- \* [\[GPU Direct\]\(../troubleshooting.html#gpu-direct\)](#)
- \* [\[GPU-to-GPU communication\]\(../troubleshooting.html#gpu-to-gpu-communication\)](#)
- \* [\[GPU-to-NIC communication\]\(../troubleshooting.html#gpu-to-nic-communication\)](#)
- \* [\[PCI Access Control Services \(ACS\)\]\(../troubleshooting.html#pci-access-control-services-ac\)](#)
- \* [\[Topology detection\]\(../troubleshooting.html#topology-detection\)](#)
- \* [\[Shared memory\]\(../troubleshooting.html#shared-memory\)](#)
- \* [\[Docker\]\(../troubleshooting.html#docker\)](#)
- \* [\[Systemd\]\(../troubleshooting.html#systemd\)](#)
- \* [\[Networking issues\]\(../troubleshooting.html#networking-issues\)](#)
- \* [\[IP Network Interfaces\]\(../troubleshooting.html#ip-network-interfaces\)](#)
- \* [\[IP Ports\]\(../troubleshooting.html#ip-ports\)](#)
- \* [\[InfiniBand\]\(../troubleshooting.html#infiniband\)](#)
- \* [\[RDMA over Converged Ethernet \(RoCE\)\]\(../troubleshooting.html#rdma-over-converged-ethernet-roce\)](#)

[\\_\\_\[NCCL\]\(../index.html\)](#)

- \* [\[Docs\]\(../index.html\)](#) »
- \* [\[Using NCCL\]\(../usage.html\)](#) »
- \* [Group Calls](#)
- \* [\[ View page source\]\(../\\_sources/usage/groups.rst.txt\)](#)

\* \* \*

# Group Calls

Group functions (`ncclGroupStart/ncclGroupEnd`) can be used to merge multiple calls into one. This is needed for three purposes: managing multiple GPUs from one thread (to avoid deadlocks), aggregating communication operations to improve performance, or merging multiple send/receive point-to-point operations (see [\[Point-to-point communication\]\(p2p.html#point-to-point\)](#) section). All three usages can be combined together, with one exception : calls to `[`ncclCommInitRank()`](../api/comms.html#c.ncclCommInitRank"ncclCommInitRank")` cannot be merged with others.

## ## Management Of Multiple GPUs From One Thread¶

When a single thread is managing multiple devices, group semantics must be used. This is because every NCCL call may have to block, waiting for other threads/ranks to arrive, before effectively posting the NCCL operation on the given stream. Hence, a simple loop on multiple devices like shown below could block on the first call waiting for the other ones:

```
for (int i=0; i<nLocalDevs; i++) {  
    ncclAllReduce(..., comm[i], stream[i]);  
}
```

To define that these calls are part of the same collective operation, `ncclGroupStart` and `ncclGroupEnd` should be used:

```
ncclGroupStart();  
  
for (int i=0; i<nLocalDevs; i++) {  
    ncclAllReduce(..., comm[i], stream[i]);  
}  
  
ncclGroupEnd();
```

This will tell NCCL to treat all calls between `ncclGroupStart` and `ncclGroupEnd` as a single call to many devices.

Caution: When called inside a group, stream operations (like `ncclAllReduce`) can return without having enqueued the operation on the stream. Stream operations like `cudaStreamSynchronize` can therefore be called only after `ncclGroupEnd` returns.

Group calls must also be used to create a communicator when one thread manages more than one device:

```
ncclGroupStart();  
  
for (int i=0; i<nLocalDevs; i++) {  
    cudaSetDevice(device[i]);  
  
    ncclCommInitRank(comms+i, nranks, commId, rank[i]);  
}
```

```
}  
  
ncclGroupEnd();
```

Note: Contrary to NCCL 1.x, there is no need to set the CUDA device before every NCCL communication call within a group, but it is still needed when calling `ncclCommInitRank` within a group.

Related links:

- \* [`ncclGroupStart()`](../api/group.html#c.ncclGroupStart "ncclGroupStart")
- \* [`ncclGroupEnd()`](../api/group.html#c.ncclGroupEnd "ncclGroupEnd")

## ## Aggregated Operations (2.2 and later)¶

The group semantics can also be used to have multiple collective operations performed within a single NCCL launch. This is useful for reducing the launch overhead, in other words, latency, as it only occurs once for multiple operations. Init functions cannot be aggregated with other init functions, nor with communication functions.

Aggregation of collective operations can be done simply by having multiple calls to NCCL within a `ncclGroupStart` / `ncclGroupEnd` section.

In the following example, we launch one broadcast and two `allReduce` operations together as a single NCCL launch.



```

ncclGroupStart();

ncclBroadcast(sendbuff1, recvbuff1, count1, datatype, root, comm, stream);

ncclAllReduce(sendbuff2, recvbuff2, count2, datatype, comm, stream);

ncclAllReduce(sendbuff3, recvbuff3, count3, datatype, comm, stream);

ncclGroupEnd();

```

It is permitted to combine aggregation with multi-GPU launch and use different communicators in a group launch as shown in the Management Of Multiple GPUs From One Thread topic. When combining multi-GPU launch and aggregation, ncclGroupStart and ncclGroupEnd can be either used once or at each level. The following example groups the allReduce operations from different layers and on multiple CUDA devices :

```

ncclGroupStart();

for (int i=0; i<nlayers; i++) {

    ncclGroupStart();

    for (int g=0; g<ngpus; g++) {

        ncclAllReduce(sendbuffs[g]+offsets[i], recvbuffs[g]+offsets[i], counts[i], datatype[i], comms[g],
streams[g]);

    }

    ncclGroupEnd();

}

```

```
ncclGroupEnd();
```

Note: The NCCL operation will only be started as a whole during the last call to `ncclGroupEnd`. The `ncclGroupStart` and `ncclGroupEnd` calls within the for loop are not necessary and do nothing.

Related links:

- \* [`ncclGroupStart()`](../api/group.html#c.ncclGroupStart "ncclGroupStart")
- \* [`ncclGroupEnd()`](../api/group.html#c.ncclGroupEnd "ncclGroupEnd")

## ## Nonblocking Group Operation¶

If a communicator is marked as nonblocking through `ncclCommInitRankConfig`, the group functions become asynchronous correspondingly. In this case, if users issue multiple NCCL operations in one group, returning from `ncclGroupEnd()` might not mean the NCCL communication kernels have been issued to CUDA streams. If `ncclGroupEnd()` returns `ncclSuccess`, it means NCCL kernels have been issued to streams; if it returns `ncclInProgress`, it means NCCL kernels are being issued to streams in the background. It is users'™ responsibility to make sure the state of the communicator changes into `ncclSuccess` before calling related CUDA calls (e.g. `cudaStreamSynchronize`):

```
ncclGroupStart();
```

```

for (int g=0; g<ngpus; g++) {

    ncclAllReduce(sendbuffs[g]+offsets[i], recvbuffs[g]+offsets[i], counts[i], datatype[i], comms[g],
streams[g]);

}

ret = ncclGroupEnd();

if (ret == ncclInProgress) {

    for (int g=0; g<ngpus; g++) {

        do {

            ncclCommGetAsyncError(comms[g], &state);

        } while (state == ncclInProgress);

    }

} else if (ret == ncclSuccess) {

    /* Successfully issued */

    printf("NCCL kernel issue succeeded\n");

} else {

    /* Errors happen */

    reportErrorAndRestart();

}

for (int g=0; g<ngpus; g++) {

    cudaStreamSynchronize(streams[g]);

}

```

Related links:

\* [\[`ncclCommInitRankConfig\(\)`\]\(../api/comms.html#c.ncclCommInitRankConfig](#)

"ncclCommInitRankConfig")

\* [\[ ncclCommGetAsyncError\(\) \]](#)(../api/comms.html#c.ncclCommGetAsyncError

"ncclCommGetAsyncError")

[\[Next \]](#)(p2p.html "Point-to-point communication") [\[ Previous\]](#)(streams.html

"CUDA Stream Semantics")

\* \* \*

(C) Copyright 2020, NVIDIA Corporation

Built with [\[Sphinx\]](#)(http://sphinx-doc.org/) using a

[\[theme\]](#)(https://github.com/rtd/sphinx\_rtd\_theme) provided by [\[Read the](#)

[Docs\]](#)(https://readthedocs.org).