

[NCCL](../index.html)

[2.25](https://docs.nvidia.com/deeplearning/sdk/nccl-archived/index.html)

- * [Overview of NCCL](../overview.html)

- * [Setup](../setup.html)

- * [Using NCCL](../usage.html)

- * [Creating a Communicator](communicators.html)

- * [Creating a communicator with options](communicators.html#creating-a-communicator-with-options)

- * [Creating a communicator using multiple ncclUniqueIds](communicators.html#creating-a-communicator-using-multiple-nccluniqueids)

- * [Creating more communicators](communicators.html#creating-more-communicators)

- * [Using multiple NCCL communicators concurrently](communicators.html#using-multiple-nccl-communicators-concurrently)

- * [Finalizing a communicator](communicators.html#finalizing-a-communicator)

- * [Destroying a communicator](communicators.html#destroying-a-communicator)

- * [Error handling and communicator abort](communicators.html#error-handling-and-communicator-abort)

- * [Asynchronous errors and error handling](communicators.html#asynchronous-errors-and-error-handling)

- * [Fault Tolerance](communicators.html#fault-tolerance)

- * [Collective Operations](collectives.html)

- * [AllReduce](collectives.html#allreduce)

- * [Broadcast](collectives.html#broadcast)

- * [Reduce](collectives.html#reduce)

- * [AllGather](collectives.html#allgather)

- * [\[ReduceScatter\]\(collectives.html#reducescatter\)](#)
- * [\[Data Pointers\]\(data.html\)](#)
- * [\[CUDA Stream Semantics\]\(streams.html\)](#)
 - * [\[Mixing Multiple Streams within the same ncclGroupStart/End\(\) group\]\(streams.html#mixing-multiple-streams-within-the-same-ncclgroupstart-end-group\)](#)
- * [\[Group Calls\]\(groups.html\)](#)
 - * [\[Management Of Multiple GPUs From One Thread\]\(groups.html#management-of-multiple-gpus-from-one-thread\)](#)
- * [\[Aggregated Operations \(2.2 and later\)\]\(groups.html#aggregated-operations-2-2-and-later\)](#)
- * [\[Nonblocking Group Operation\]\(groups.html#nonblocking-group-operation\)](#)
- * Point-to-point communication
 - * Sendrecv
 - * One-to-all (scatter)
 - * All-to-one (gather)
 - * All-to-all
 - * Neighbor exchange
- * [\[Thread Safety\]\(threadsafety.html\)](#)
- * [\[In-place Operations\]\(inplace.html\)](#)
- * [\[Using NCCL with CUDA Graphs\]\(cudagraph.html\)](#)
- * [\[User Buffer Registration\]\(bufferreg.html\)](#)
 - * [\[NVLink Sharp Buffer Registration\]\(bufferreg.html#nvlink-sharp-buffer-registration\)](#)
 - * [\[IB Sharp Buffer Registration\]\(bufferreg.html#ib-sharp-buffer-registration\)](#)
 - * [\[General Buffer Registration\]\(bufferreg.html#general-buffer-registration\)](#)
 - * [\[Memory Allocator\]\(bufferreg.html#memory-allocator\)](#)
- * [\[NCCL API\]\(../api.html\)](#)
 - * [\[Communicator Creation and Management Functions\]\(../api/comms.html\)](#)
 - * [\[ncclGetLastError\]\(../api/comms.html#ncclgetlasterror\)](#)

- * [\[ncclGetErrorString\]\(../api/comms.html#ncclgeterrorstring\)](#)
- * [\[ncclGetVersion\]\(../api/comms.html#ncclgetversion\)](#)
- * [\[ncclGetUniqueId\]\(../api/comms.html#ncclgetuniqueid\)](#)
- * [\[ncclCommInitRank\]\(../api/comms.html#ncclcomminitrank\)](#)
- * [\[ncclCommInitAll\]\(../api/comms.html#ncclcomminitall\)](#)
- * [\[ncclCommInitRankConfig\]\(../api/comms.html#ncclcomminitrankconfig\)](#)
- * [\[ncclCommInitRankScalable\]\(../api/comms.html#ncclcomminitrankscalable\)](#)
- * [\[ncclCommSplit\]\(../api/comms.html#ncclcommsplit\)](#)
- * [\[ncclCommFinalize\]\(../api/comms.html#ncclcommfinalize\)](#)
- * [\[ncclCommDestroy\]\(../api/comms.html#ncclcommdestroy\)](#)
- * [\[ncclCommAbort\]\(../api/comms.html#ncclcommabort\)](#)
- * [\[ncclCommGetAsyncError\]\(../api/comms.html#ncclcommgetasyncerror\)](#)
- * [\[ncclCommCount\]\(../api/comms.html#ncclcommcount\)](#)
- * [\[ncclCommCuDevice\]\(../api/comms.html#ncclcommcudevice\)](#)
- * [\[ncclCommUserRank\]\(../api/comms.html#ncclcommuserrank\)](#)
- * [\[ncclCommRegister\]\(../api/comms.html#ncclcommregister\)](#)
- * [\[ncclCommDeregister\]\(../api/comms.html#ncclcommderegister\)](#)
- * [\[ncclMemAlloc\]\(../api/comms.html#ncclmemalloc\)](#)
- * [\[ncclMemFree\]\(../api/comms.html#ncclmemfree\)](#)
- * [\[Collective Communication Functions\]\(../api/colls.html\)](#)
 - * [\[ncclAllReduce\]\(../api/colls.html#ncclallreduce\)](#)
 - * [\[ncclBroadcast\]\(../api/colls.html#ncclbroadcast\)](#)
 - * [\[ncclReduce\]\(../api/colls.html#ncclreduce\)](#)
 - * [\[ncclAllGather\]\(../api/colls.html#ncclallgather\)](#)
 - * [\[ncclReduceScatter\]\(../api/colls.html#ncclreducescatter\)](#)
- * [\[Group Calls\]\(../api/group.html\)](#)
 - * [\[ncclGroupStart\]\(../api/group.html#ncclgroupstart\)](#)

- * [\[ncclGroupEnd\]\(../api/group.html#ncclgroupend\)](#)
- * [\[ncclGroupSimulateEnd\]\(../api/group.html#ncclgroupsimulateend\)](#)
- * [\[Point To Point Communication Functions\]\(../api/p2p.html\)](#)
- * [\[ncclSend\]\(../api/p2p.html#ncclsend\)](#)
- * [\[ncclRecv\]\(../api/p2p.html#ncclrecv\)](#)
- * [\[Types\]\(../api/types.html\)](#)
- * [\[ncclComm_t\]\(../api/types.html#ncclcomm-t\)](#)
- * [\[ncclResult_t\]\(../api/types.html#ncclresult-t\)](#)
- * [\[ncclDataType_t\]\(../api/types.html#nccldatatype-t\)](#)
- * [\[ncclRedOp_t\]\(../api/types.html#ncclredop-t\)](#)
- * [\[ncclScalarResidence_t\]\(../api/types.html#ncclscalarresidence-t\)](#)
- * [\[ncclConfig_t\]\(../api/types.html#ncclconfig-t\)](#)
- * [\[ncclSimInfo_t\]\(../api/types.html#ncclsiminfo-t\)](#)
- * [\[User Defined Reduction Operators\]\(../api/ops.html\)](#)
- * [\[ncclRedOpCreatePreMulSum\]\(../api/ops.html#ncclredopcreatepremulsum\)](#)
- * [\[ncclRedOpDestroy\]\(../api/ops.html#ncclredopdestroy\)](#)
- * [\[Migrating from NCCL 1 to NCCL 2\]\(../nccl1.html\)](#)
- * [\[Initialization\]\(../nccl1.html#initialization\)](#)
- * [\[Communication\]\(../nccl1.html#communication\)](#)
- * [\[Counts\]\(../nccl1.html#counts\)](#)
- * [\[In-place usage for AllGather and ReduceScatter\]\(../nccl1.html#in-place-usage-for-allgather-and-reducescatter\)](#)
- * [\[AllGather arguments order\]\(../nccl1.html#allgather-arguments-order\)](#)
- * [\[Datatypes\]\(../nccl1.html#datatypes\)](#)
- * [\[Error codes\]\(../nccl1.html#error-codes\)](#)
- * [\[Examples\]\(../examples.html\)](#)
- * [\[Communicator Creation and Destruction\]](#)

Examples](../examples.html#communicator-creation-and-destruction-examples)

- * [Example 1: Single Process, Single Thread, Multiple Devices](../examples.html#example-1-single-process-single-thread-multiple-devices)

- * [Example 2: One Device per Process or Thread](../examples.html#example-2-one-device-per-process-or-thread)

- * [Example 3: Multiple Devices per Thread](../examples.html#example-3-multiple-devices-per-thread)

- * [Example 4: Multiple communicators per device](../examples.html#example-4-multiple-communicators-per-device)

- * [Communication Examples](../examples.html#communication-examples)

- * [Example 1: One Device per Process or Thread](../examples.html#example-1-one-device-per-process-or-thread)

- * [Example 2: Multiple Devices per Thread](../examples.html#example-2-multiple-devices-per-thread)

- * [NCCL and MPI](../mpi.html)

- * [API](../mpi.html#api)

- * [Using multiple devices per process](../mpi.html#using-multiple-devices-per-process)

- * [ReduceScatter operation](../mpi.html#reducescatter-operation)

- * [Send and Receive counts](../mpi.html#send-and-receive-counts)

- * [Other collectives and point-to-point operations](../mpi.html#other-collectives-and-point-to-point-operations)

- * [In-place operations](../mpi.html#in-place-operations)

- * [Using NCCL within an MPI Program](../mpi.html#using-nccl-within-an-mpi-program)

- * [MPI Progress](../mpi.html#mpi-progress)

- * [Inter-GPU Communication with CUDA-aware MPI](../mpi.html#inter-gpu-communication-with-cuda-aware-mpi)

- * [Environment Variables](../env.html)

* [System configuration](../env.html#system-configuration)

* [NCCL_SOCKET_IFNAME](../env.html#nccl-socket-ifname)

* [Values accepted](../env.html#values-accepted)

* [NCCL_SOCKET_FAMILY](../env.html#nccl-socket-family)

* [Values accepted](../env.html#id2)

* [NCCL_SOCKET_RETRY_CNT](../env.html#nccl-socket-retry-cnt)

* [Values accepted](../env.html#id3)

* [NCCL_SOCKET_RETRY_SLEEP_MSEC](../env.html#nccl-socket-retry-sleep-msec)

* [Values accepted](../env.html#id4)

* [NCCL_SOCKET_NTHREADS](../env.html#nccl-socket-nthreads)

* [Values accepted](../env.html#id5)

* [NCCL_NSOCKS_PERTHREAD](../env.html#nccl-nsocks-perthread)

* [Values accepted](../env.html#id6)

* [NCCL_CROSS_NIC](../env.html#nccl-cross-nic)

* [Values accepted](../env.html#id7)

* [NCCL_IB_HCA](../env.html#nccl-ib-hca)

* [Values accepted](../env.html#id8)

* [NCCL_IB_TIMEOUT](../env.html#nccl-ib-timeout)

* [Values accepted](../env.html#id9)

* [NCCL_IB_RETRY_CNT](../env.html#nccl-ib-retry-cnt)

* [Values accepted](../env.html#id10)

* [NCCL_IB_GID_INDEX](../env.html#nccl-ib-gid-index)

* [Values accepted](../env.html#id11)

* [NCCL_IB_ADDR_FAMILY](../env.html#nccl-ib-addr-family)

* [Values accepted](../env.html#id12)

* [NCCL_IB_ADDR_RANGE](../env.html#nccl-ib-addr-range)

* [Values accepted](../env.html#id13)

* [NCCL_IB_ROCE_VERSION_NUM](../env.html#nccl-ib-roce-version-num)
* [Values accepted](../env.html#id14)

* [NCCL_IB_SL](../env.html#nccl-ib-sl)
* [Values accepted](../env.html#id15)

* [NCCL_IB_TC](../env.html#nccl-ib-tc)
* [Values accepted](../env.html#id16)

* [NCCL_IB_FIFO_TC](../env.html#nccl-ib-fifo-tc)
* [Values accepted](../env.html#id17)

* [NCCL_IB_RETURN_ASYNC_EVENTS](../env.html#nccl-ib-return-async-events)
* [Values accepted](../env.html#id18)

* [NCCL_OOB_NET_ENABLE](../env.html#nccl-oob-net-enable)
* [Values accepted](../env.html#id19)

* [NCCL_OOB_NET_IFNAME](../env.html#nccl-oob-net-ifname)
* [Values accepted](../env.html#id20)

* [NCCL_UID_STAGGER_THRESHOLD](../env.html#nccl-uid-stagger-threshold)
* [Values accepted](../env.html#id21)

* [NCCL_UID_STAGGER_RATE](../env.html#nccl-uid-stagger-rate)
* [Values accepted](../env.html#id22)

* [NCCL_NET](../env.html#nccl-net)
* [Values accepted](../env.html#id23)

* [NCCL_NET_PLUGIN](../env.html#nccl-net-plugin)
* [Values accepted](../env.html#id24)

* [NCCL_TUNER_PLUGIN](../env.html#nccl-tuner-plugin)
* [Values accepted](../env.html#id25)

* [NCCL_PROFILER_PLUGIN](../env.html#nccl-profiler-plugin)
* [Values accepted](../env.html#id26)

* [NCCL_IGNORE_CPU_AFFINITY](../env.html#nccl-ignore-cpu-affinity)

- * [Values accepted](../env.html#id27)
- * [NCCL_CONF_FILE](../env.html#nccl-conf-file)
- * [Values accepted](../env.html#id28)
- * [NCCL_DEBUG](../env.html#nccl-debug)
- * [Values accepted](../env.html#id30)
- * [NCCL_DEBUG_FILE](../env.html#nccl-debug-file)
- * [Values accepted](../env.html#id31)
- * [NCCL_DEBUG_SUBSYS](../env.html#nccl-debug-subsys)
- * [Values accepted](../env.html#id32)
- * [NCCL_COLLNET_ENABLE](../env.html#nccl-collnet-enable)
- * [Value accepted](../env.html#value-accepted)
- * [NCCL_COLLNET_NODE_THRESHOLD](../env.html#nccl-collnet-node-threshold)
- * [Value accepted](../env.html#id33)
- * [NCCL_TOPO_FILE](../env.html#nccl-topo-file)
- * [Value accepted](../env.html#id34)
- * [NCCL_TOPO_DUMP_FILE](../env.html#nccl-topo-dump-file)
- * [Value accepted](../env.html#id35)
- * [NCCL_SET_THREAD_NAME](../env.html#nccl-set-thread-name)
- * [Value accepted](../env.html#id36)
- * [Debugging](../env.html#debugging)
- * [NCCL_P2P_DISABLE](../env.html#nccl-p2p-disable)
- * [Values accepted](../env.html#id37)
- * [NCCL_P2P_LEVEL](../env.html#nccl-p2p-level)
- * [Values accepted](../env.html#id38)
- * [Integer Values (Legacy)](../env.html#integer-values-legacy)
- * [NCCL_P2P_DIRECT_DISABLE](../env.html#nccl-p2p-direct-disable)
- * [Values accepted](../env.html#id39)

* [NCCL_SHM_DISABLE](../env.html#nccl-shm-disable)
* [Values accepted](../env.html#id40)

* [NCCL_BUFFSIZE](../env.html#nccl-buffersize)
* [Values accepted](../env.html#id41)

* [NCCL_NTHREADS](../env.html#nccl-nthreads)
* [Values accepted](../env.html#id42)

* [NCCL_MAX_NCHANNELS](../env.html#nccl-max-nchannels)
* [Values accepted](../env.html#id43)

* [NCCL_MIN_NCHANNELS](../env.html#nccl-min-nchannels)
* [Values accepted](../env.html#id44)

* [NCCL_CHECKS_DISABLE](../env.html#nccl-checks-disable)
* [Values accepted](../env.html#id45)

* [NCCL_CHECK_POINTERS](../env.html#nccl-check-pointers)
* [Values accepted](../env.html#id46)

* [NCCL_LAUNCH_MODE](../env.html#nccl-launch-mode)
* [Values accepted](../env.html#id47)

* [NCCL_IB_DISABLE](../env.html#nccl-ib-disable)
* [Values accepted](../env.html#id48)

* [NCCL_IB_AR_THRESHOLD](../env.html#nccl-ib-ar-threshold)
* [Values accepted](../env.html#id49)

* [NCCL_IB_QPS_PER_CONNECTION](../env.html#nccl-ib-qps-per-connection)
* [Values accepted](../env.html#id50)

* [NCCL_IB_SPLIT_DATA_ON_QPS](../env.html#nccl-ib-split-data-on-qps)
* [Values accepted](../env.html#id51)

* [NCCL_IB_CUDA_SUPPORT](../env.html#nccl-ib-cuda-support)
* [Values accepted](../env.html#id52)

* [NCCL_IB_PCI_RELAXED_ORDERING](../env.html#nccl-ib-pci-relaxed-ordering)

* [Values accepted](../env.html#id53)

* [NCCL_IB_ADAPTIVE_ROUTING](../env.html#nccl-ib-adaptive-routing)

* [Values accepted](../env.html#id54)

* [NCCL_IB_ECE_ENABLE](../env.html#nccl-ib-ece-enable)

* [Values accepted](../env.html#id55)

* [NCCL_MEM_SYNC_DOMAIN](../env.html#nccl-mem-sync-domain)

* [Values accepted](../env.html#id56)

* [NCCL_CUMEM_ENABLE](../env.html#nccl-cumem-enable)

* [Values accepted](../env.html#id57)

* [NCCL_CUMEM_HOST_ENABLE](../env.html#nccl-cumem-host-enable)

* [Values accepted](../env.html#id58)

* [NCCL_NET_GDR_LEVEL (formerly

NCCL_IB_GDR_LEVEL)](../env.html#nccl-net-gdr-level-formerly-nccl-ib-gdr-level)

* [Values accepted](../env.html#id59)

* [Integer Values (Legacy)](../env.html#id60)

* [NCCL_NET_GDR_READ](../env.html#nccl-net-gdr-read)

* [Values accepted](../env.html#id61)

* [NCCL_NET_SHARED_BUFFERS](../env.html#nccl-net-shared-buffers)

* [Value accepted](../env.html#id62)

* [NCCL_NET_SHARED_COMMS](../env.html#nccl-net-shared-comms)

* [Value accepted](../env.html#id63)

* [NCCL_SINGLE_RING_THRESHOLD](../env.html#nccl-single-ring-threshold)

* [Values accepted](../env.html#id64)

* [NCCL_LL_THRESHOLD](../env.html#nccl-ll-threshold)

* [Values accepted](../env.html#id65)

* [NCCL_TREE_THRESHOLD](../env.html#nccl-tree-threshold)

* [Values accepted](../env.html#id66)

* [NCCL_ALGO](../env.html#nccl-algo)

* [Values accepted](../env.html#id67)

* [NCCL_PROTO](../env.html#nccl-proto)

* [Values accepted](../env.html#id68)

* [NCCL_NVX_DISABLE](../env.html#nccl-nvx-disable)

* [Value accepted](../env.html#id69)

* [NCCL_P2P_DISABLE](../env.html#nccl-p2p-disable)

* [Value accepted](../env.html#id70)

* [NCCL_P2P_P2P_LEVEL](../env.html#nccl-p2p-p2p-level)

* [Value accepted](../env.html#id71)

* [NCCL_RUNTIME_CONNECT](../env.html#nccl-runtime-connect)

* [Value accepted](../env.html#id72)

* [NCCL_GRAPH_REGISTER](../env.html#nccl-graph-register)

* [Value accepted](../env.html#id74)

* [NCCL_LOCAL_REGISTER](../env.html#nccl-local-register)

* [Value accepted](../env.html#id75)

* [NCCL_LEGACY_CUDA_REGISTER](../env.html#nccl-legacy-cuda-register)

* [Value accepted](../env.html#id76)

* [NCCL_SET_STACK_SIZE](../env.html#nccl-set-stack-size)

* [Value accepted](../env.html#id77)

* [NCCL_GRAPH_MIXING_SUPPORT](../env.html#nccl-graph-mixing-support)

* [Value accepted](../env.html#id79)

* [NCCL_DMABUF_ENABLE](../env.html#nccl-dmabuf-enable)

* [Value accepted](../env.html#id80)

* [NCCL_P2P_NET_CHUNKSIZE](../env.html#nccl-p2p-net-chunksize)

* [Values accepted](../env.html#id81)

* [NCCL_P2P_LL_THRESHOLD](../env.html#nccl-p2p-ll-threshold)

- * [\[Values accepted\]\(../env.html#id82\)](#)
- * [\[NCCL_ALLOC_P2P_NET_LL_BUFFERS\]\(../env.html#nccl-alloc-p2p-net-ll-buffers\)](#)
- * [\[Values accepted\]\(../env.html#id83\)](#)
- * [\[NCCL_COMM_BLOCKING\]\(../env.html#nccl-comm-blocking\)](#)
- * [\[Values accepted\]\(../env.html#id84\)](#)
- * [\[NCCL_CGA_CLUSTER_SIZE\]\(../env.html#nccl-cga-cluster-size\)](#)
- * [\[Values accepted\]\(../env.html#id85\)](#)
- * [\[NCCL_MAX_CTAS\]\(../env.html#nccl-max-ctas\)](#)
- * [\[Values accepted\]\(../env.html#id86\)](#)
- * [\[NCCL_MIN_CTAS\]\(../env.html#nccl-min-ctas\)](#)
- * [\[Values accepted\]\(../env.html#id87\)](#)
- * [\[NCCL_NVLS_ENABLE\]\(../env.html#nccl-nvls-enable\)](#)
- * [\[Values accepted\]\(../env.html#id88\)](#)
- * [\[NCCL_IB_MERGE_NICS\]\(../env.html#nccl-ib-merge-nics\)](#)
- * [\[Values accepted\]\(../env.html#id89\)](#)
- * [\[NCCL_MNNVL_ENABLE\]\(../env.html#nccl-mnnvl-enable\)](#)
- * [\[Values accepted\]\(../env.html#id90\)](#)
- * [\[NCCL_RAS_ENABLE\]\(../env.html#nccl-ras-enable\)](#)
- * [\[Values accepted\]\(../env.html#id91\)](#)
- * [\[NCCL_RAS_ADDR\]\(../env.html#nccl-ras-addr\)](#)
- * [\[Values accepted\]\(../env.html#id92\)](#)
- * [\[NCCL_RAS_TIMEOUT_FACTOR\]\(../env.html#nccl-ras-timeout-factor\)](#)
- * [\[Values accepted\]\(../env.html#id93\)](#)
- * [\[Troubleshooting\]\(../troubleshooting.html\)](#)
- * [\[Errors\]\(../troubleshooting.html#errors\)](#)
- * [\[RAS\]\(../troubleshooting.html#ras\)](#)
- * [\[RAS\]\(../troubleshooting/ras.html\)](#)

- * [\[Principle of Operation\]\(../troubleshooting/ras.html#principle-of-operation\)](#)
- * [\[RAS Queries\]\(../troubleshooting/ras.html#ras-queries\)](#)
- * [\[Sample Output\]\(../troubleshooting/ras.html#sample-output\)](#)
- * [\[GPU Direct\]\(../troubleshooting.html#gpu-direct\)](#)
- * [\[GPU-to-GPU communication\]\(../troubleshooting.html#gpu-to-gpu-communication\)](#)
- * [\[GPU-to-NIC communication\]\(../troubleshooting.html#gpu-to-nic-communication\)](#)
- * [\[PCI Access Control Services \(ACS\)\]\(../troubleshooting.html#pci-access-control-services-ac\)](#)
- * [\[Topology detection\]\(../troubleshooting.html#topology-detection\)](#)
- * [\[Shared memory\]\(../troubleshooting.html#shared-memory\)](#)
- * [\[Docker\]\(../troubleshooting.html#docker\)](#)
- * [\[Systemd\]\(../troubleshooting.html#systemd\)](#)
- * [\[Networking issues\]\(../troubleshooting.html#networking-issues\)](#)
- * [\[IP Network Interfaces\]\(../troubleshooting.html#ip-network-interfaces\)](#)
- * [\[IP Ports\]\(../troubleshooting.html#ip-ports\)](#)
- * [\[InfiniBand\]\(../troubleshooting.html#infiniband\)](#)
- * [\[RDMA over Converged Ethernet \(RoCE\)\]\(../troubleshooting.html#rdma-over-converged-ethernet-roce\)](#)

[__\[NCCL\]\(../index.html\)](#)

- * [\[Docs\]\(../index.html\)](#) »
- * [\[Using NCCL\]\(../usage.html\)](#) »
- * [Point-to-point communication](#)
- * [\[View page source\]\(../_sources/usage/p2p.rst.txt\)](#)

* * *

Point-to-point communication

(Since NCCL 2.7) Point-to-point communication can be used to express any communication pattern between ranks. Any point-to-point communication needs two NCCL calls : a call to `ncclSend()` ([./api/p2p.html#c.ncclSend](https://nvidia.github.io/nccl/2.7.0/api/p2p.html#c.ncclSend) "ncclSend") on one rank and a corresponding `ncclRecv()` ([./api/p2p.html#c.ncclRecv](https://nvidia.github.io/nccl/2.7.0/api/p2p.html#c.ncclRecv) "ncclRecv") on the other rank, with the same count and data type.

Multiple calls to `ncclSend()` ([./api/p2p.html#c.ncclSend](https://nvidia.github.io/nccl/2.7.0/api/p2p.html#c.ncclSend) "ncclSend") and `ncclRecv()` ([./api/p2p.html#c.ncclRecv](https://nvidia.github.io/nccl/2.7.0/api/p2p.html#c.ncclRecv) "ncclRecv") targeting different peers can be fused together with `ncclGroupStart()` ([./api/group.html#c.ncclGroupStart](https://nvidia.github.io/nccl/2.7.0/api/group.html#c.ncclGroupStart) "ncclGroupStart") and `ncclGroupEnd()` ([./api/group.html#c.ncclGroupEnd](https://nvidia.github.io/nccl/2.7.0/api/group.html#c.ncclGroupEnd) "ncclGroupEnd") to form more complex communication patterns such as one-to-all (scatter), all-to-one (gather), all-to-all or communication with neighbors in an N-dimensional space.

Point-to-point calls within a group will be blocking until that group of calls completes, but calls within a group can be seen as progressing independently, hence should never block each other. It is therefore important to merge calls that need to progress concurrently to avoid deadlocks.

Below are a few examples of classic point-to-point communication patterns used by parallel applications. NCCL semantics allow for all variants with different sizes, datatypes, and buffers, per rank.

Sendrecv

In MPI terms, a sendrecv operation is when two ranks exchange data, both sending and receiving at the same time. This can be done by merging both ncclSend and ncclRecv calls into one :

```
ncclGroupStart();  
  
ncclSend(sendbuff, sendcount, sendtype, peer, comm, stream);  
  
ncclRecv(recvbuff, recvcount, recvtype, peer, comm, stream);  
  
ncclGroupEnd();
```

One-to-all (scatter)

A one-to-all operation from a `root` rank can be expressed by merging all send and receive operations in a group :

```
ncclGroupStart();  
  
if (rank == root) {  
    for (int r=0; r<n ranks; r++)  
        ncclSend(sendbuff[r], size, type, r, comm, stream);  
}  
  
ncclRecv(recvbuff, size, type, root, comm, stream);
```

```
ncclGroupEnd();
```

```
## All-to-one (gather)¶
```

Similarly, an all-to-one operations to a `root` rank would be implemented this way :

```
ncclGroupStart();  
if (rank == root) {  
    for (int r=0; r<n ranks; r++)  
        ncclRecv(recvbuff[r], size, type, r, comm, stream);  
}  
ncclSend(sendbuff, size, type, root, comm, stream);  
ncclGroupEnd();
```

```
## All-to-all¶
```

An all-to-all operation would be a merged loop of send/recv operations to/from all peers :

```
ncclGroupStart();
```



```

for (int r=0; r<n ranks; r++) {

    ncclSend(sendbuff[r], sendcount, sendtype, r, comm, stream);

    ncclRecv(recvbuff[r], recvcount, recvtype, r, comm, stream);

}

ncclGroupEnd();

```

Neighbor exchange

Finally, exchanging data with neighbors in an N-dimensions space could be done with :

```

ncclGroupStart();

for (int d=0; d<ndims; d++) {

    ncclSend(sendbuff[d], sendcount, sendtype, next[d], comm, stream);

    ncclRecv(recvbuff[d], recvcount, recvtype, prev[d], comm, stream);

}

ncclGroupEnd();

```

[\[Next \]\(threadsafety.html "Thread Safety"\)](#) [\[Previous\]\(groups.html "Group Calls"\)](#)

* * *

(C) Copyright 2020, NVIDIA Corporation

Built with [Sphinx](<http://sphinx-doc.org/>) using a

[theme](https://github.com/rtfd/sphinx_rtd_theme) provided by [Read the

Docs](<https://readthedocs.org>).