 Hugging Face

- * [Models](#)
- * [Datasets](#)
- * [Spaces](#)
- * [Posts](#)
- * [Docs](#)
- * [Enterprise](#)
- * [Pricing](#)
- * * * * *

- * [Log In](#)
- * [Sign Up](#)

#



/

[e5-mistral-7b-instruct](#)

like 486

[Feature Extraction](/models?pipeline_tag=feature-extraction)[sentence-transformers](/models?library=sentence-transformers)[PyTorch](/models?library=pytorch)[Safetensors](/models?library=safetensors)[Transformers](/models?library=transformers)[English](/models?language=en)[mistral](/models?other=mistral)[mteb](/models?other=mteb)[Eval Results](/models?other=model-index)[text-generation-inference](/models?other=text-generation-inference)[text-embeddings-inference](/models?other=text-embeddings-inference)[Inference Endpoints](/models?other=endpoints_compatible)

arxiv: 2401.00368

arxiv: 2104.08663

arxiv: 2210.07316

License: mit

[Model card](/intfloat/e5-mistral-7b-instruct)[Files Files and versions](/intfloat/e5-mistral-7b-instruct/tree/main)[Community 47](/intfloat/e5-mistral-7b-instruct/discussions)

Train

Deploy

Use this model

[\[New discussion \]\(/intfloat/e5-mistral-7b-instruct/discussions/new\)](/intfloat/e5-mistral-7b-instruct/discussions/new)

New pull request

Resources

* [\[PR & discussions documentation \]\(/docs/hub/repositories-pull-requests-discussions\)](/docs/hub/repositories-pull-requests-discussions)

* [\[Code of Conduct \]\(/code-of-conduct\)](/code-of-conduct)

* [\[Hub documentation \]\(/docs/hub\)](/docs/hub)

* All


* Discussions

* Pull requests

[\[View closed \(13\) \]\(/intfloat/e5-mistral-7b-instruct/discussions?status=open\)](/intfloat/e5-mistral-7b-instruct/discussions?status=open)


[\[Question about MTEB Evaluation and max_seq_length Settings for e5-mistral-7b-instruct \]\(/intfloat/e5-mistral-7b-instruct/discussions/47\)](/intfloat/e5-mistral-7b-instruct/discussions/47)

#47 opened about 2 months ago by

 [george31 \]\(/george31\)](/george31)

[\[Is it possible to change the dimentions of the embeddings it generates? \]\(/intfloat/e5-mistral-7b-instruct/discussions/46\)](/intfloat/e5-mistral-7b-instruct/discussions/46)


#46 opened 4 months ago by

 srj

[Adding `safetensors` variant of this model

](/intfloat/e5-mistral-7b-instruct/discussions/45)


#45 opened 8 months ago by

 jakeyap

[Best Practices for Fine-Tuning Models on Multi-Hop Datasets?

](/intfloat/e5-mistral-7b-instruct/discussions/43)


#43 opened 8 months ago by

 leonshub

[How to set instruct for MRACL dataset? 1

](/intfloat/e5-mistral-7b-instruct/discussions/41)

#41 opened 10 months ago by

 Mengyao00

[CUDA Memory error when using sentence transformers using Tesla

V100-PCI-E-32GB 2](/intfloat/e5-mistral-7b-instruct/discussions/40)

#40 opened 10 months ago by

 al-h

[How did you construct homogeneous batch if there is not enough sample

to make single batch?](/intfloat/e5-mistral-7b-instruct/discussions/38)

#38 opened 10 months ago by

[!](/avatars/aa12efb67bcbf412d50fb9318f1db4b3.svg) lightapple](/lightapple)

[Can anyone help on this phenomenon? Fine tuning intfloat /

e5-mistral-7b-instruct got nan loss and all NaN intermediate model outputs

](/intfloat/e5-mistral-7b-instruct/discussions/37)

#37 opened 10 months ago by

[!](/avatars/5167412cd446d022faebc2df4f94447e.svg) jhyapple](/jhyapple)

[Use in sentence transformers

](/intfloat/e5-mistral-7b-instruct/discussions/35)

#35 opened 11 months ago by

[!](/avatars/4db9531e9d9d7117e253216130579386.svg) woofadu](/woofadu)

[Understanding memory consumption during inference 2

](/intfloat/e5-mistral-7b-instruct/discussions/34)


#34 opened 11 months ago by

[!](/avatars/5a43c0cf06d692774295158c7fc5aa08.svg) andrew-kirfman](/andrew-kirfman)

[Questions about the training data 4

](/intfloat/e5-mistral-7b-instruct/discussions/33)

#33 opened 11 months ago by

 xxxcliu

[How much GPU memory is required for 32k context embedding? 1


](/intfloat/e5-mistral-7b-instruct/discussions/32)

#32 opened 11 months ago by

 Labmem009

[Quantization 1](/intfloat/e5-mistral-7b-instruct/discussions/30)


#30 opened 12 months ago by

 mrgiraffe

[Train loss for reproduction

](/intfloat/e5-mistral-7b-instruct/discussions/29)


#29 opened 12 months ago by

 yjkimstats

[Is there a way to do fine tuning using STS datasets?


](/intfloat/e5-mistral-7b-instruct/discussions/27)

#27 opened about 1 year ago by

 ijkim

[Instruction / Query document embedding question. 2

](/intfloat/e5-mistral-7b-instruct/discussions/26)

#26 opened about 1 year ago by  (https://cdn-


avatars.huggingface.co/v1/production/uploads/630f2745982455e61cc5fb1d/Pi7Y2p5R2UI5tMGCIX5vy.jpeg)

eeek](/eeek)

[KeyError: 'mistral' Application startup failed during deployment on

Inference Endpoints 1](/intfloat/e5-mistral-7b-instruct/discussions/24)

#24 opened about 1 year ago by

 (/avatars/2e5daec5b810487a6bf91aac75293885.svg) hoangsx](/hoangsx)

[best way to deploy this as an API endpoint. 12

](/intfloat/e5-mistral-7b-instruct/discussions/21)

#21 opened about 1 year ago by  (https://cdn-


avatars.huggingface.co/v1/production/uploads/6454631806728ff79a3829e8/lm1TjNIXt7J5Am4aW6FZ-.jpeg)

aigeek0x0](/aigeek0x0)

[Error using multiple GPUs 4

](/intfloat/e5-mistral-7b-instruct/discussions/20)

#20 opened about 1 year ago by

 (/avatars/b78bbc219a9a4fbd0aa38b86c522887.svg) JesusUned](/JesusUned)

[Can you tell me the detail about negative sampling for benchmark data?

2](/intfloat/e5-mistral-7b-instruct/discussions/19)

#19 opened about 1 year ago by

[!](/avatars/aa12efb67bcbf412d50fb9318f1db4b3.svg) lightapple](/lightapple)

[Deployment fails (weight model.embed_tokens.weight does not exist) to

Inference Endpoint 5](/intfloat/e5-mistral-7b-instruct/discussions/18)

#18 opened about 1 year ago by

[!](/avatars/9a7ad36176d9ffe7bff382f3dd0e9d05.svg) ts314](/ts314)

[how to extract general text embedding

8](/intfloat/e5-mistral-7b-instruct/discussions/17)

#17 opened about 1 year ago by

[!](/avatars/6979d2bc37ed74822bcb115dc2d1b43f.svg)

daisyyayueyue](/daisyyayueyue)

[Is it possible to convert the model to ONNX? 1

](/intfloat/e5-mistral-7b-instruct/discussions/16)

#16 opened about 1 year ago by

[!](/avatars/d05dfbb419636a22a12abd5a15e8c875.svg)

serdarcaglar](/serdarcaglar)

[About max sequence length

2](/intfloat/e5-mistral-7b-instruct/discussions/14)

#14 opened about 1 year ago by

[!](/avatars/254021597e6d0503f5c81dae7aa5a221.svg) jorisfu](/jorisfu)

[Knowledge distillation into smaller model 2

](/intfloat/e5-mistral-7b-instruct/discussions/13)

#13 opened about 1 year ago by [!](https://cdn-

avatars.huggingface.co/v1/production/uploads/6317233cc92fd6fee317e030/cJHSvvimr1kqgQfHOjO5n.png)

tomaarsen](/tomaarsen)

[RuntimeError: weight model.embed_tokens.weight does not exist 6

](/intfloat/e5-mistral-7b-instruct/discussions/12)

#12 opened about 1 year ago by [!](https://cdn-

avatars.huggingface.co/v1/production/uploads/64d206b162a5593f17c7c72e/undwgFEhwIQ-IIRj9K35A.png) ivankeller](/ivankeller)

[When should I use this instruction for RAG retrival tasks? 5

](/intfloat/e5-mistral-7b-instruct/discussions/11)

#11 opened about 1 year ago by

[!](/avatars/ce61181898e298cbfd69ecfedc8082e9.svg) kk3dmax](/kk3dmax)

[Hardware Requirements 4

](/intfloat/e5-mistral-7b-instruct/discussions/10)

#10 opened about 1 year ago by

[!](/avatars/de6625da6088b8f29d5713fb407e4197.svg) Ravnoor1](/Ravnoor1)

[Can you training a multilingual e5 based on a multilingual LLM? 3

](/intfloat/e5-mistral-7b-instruct/discussions/9)

#9 opened about 1 year ago by

[!](/avatars/1bb0ca22bd97a9e8996e31e7b02853e6.svg) hantian](/hantian)

[symmetric tasks prompts 1

](/intfloat/e5-mistral-7b-instruct/discussions/8)

#8 opened about 1 year ago by

[!](/avatars/644e6c61eb4b7ba675cc745ced7bcd50.svg) mayjacky](/mayjacky)

[Any plan to release fine-tuning scripts? 4

](/intfloat/e5-mistral-7b-instruct/discussions/7)

#7 opened about 1 year ago by

[!](/avatars/9a5787939ecc7dfa0cd56b1d6108d466.svg) Mengyao00](/Mengyao00)

[lora config 1](/intfloat/e5-mistral-7b-instruct/discussions/6)

#6 opened about 1 year ago by [!](https://cdn-

avatars.huggingface.co/v1/production/uploads/635cc29de7aef2358a9b03ee/SVHL_mTCiOfmBamz
Sucb0.jpeg)
SeanLee97](/SeanLee97)

[Any plan to release training data? 3
](/intfloat/e5-mistral-7b-instruct/discussions/3)

#3 opened about 1 year ago by
[!](/avatars/9cf00a775b80230aa32ef68ea7377833.svg) memray](/memray)

[MLX example 3](/intfloat/e5-mistral-7b-instruct/discussions/1)

#1 opened about 1 year ago by
[!](/avatars/66071cb9b1cb4bae3c0acc2d4ca55450.svg)
paulmaksimovich](/paulmaksimovich)

System theme

Company

[TOS](/terms-of-service) [Privacy](/privacy) [About](/huggingface)
[Jobs](https://apply.workable.com/huggingface/) [](/)

Website

[Models](/models) [Datasets](/datasets) [Spaces](/spaces) [Pricing](/pricing)
[Docs](/docs)

