

(deployment-k8s)=

## # Using Kubernetes

Using Kubernetes to deploy vLLM is a scalable and efficient way to serve machine learning models. This guide will walk you through the process of deploying vLLM with Kubernetes, including the necessary prerequisites, steps for deployment, and testing.

## ## Prerequisites

Before you begin, ensure that you have the following:

- A running Kubernetes cluster
- NVIDIA Kubernetes Device Plugin (`k8s-device-plugin`): This can be found at [`https://github.com/NVIDIA/k8s-device-plugin/`](https://github.com/NVIDIA/k8s-device-plugin/)
- Available GPU resources in your cluster

## ## Deployment Steps

### 1. Create a PVC, Secret and Deployment for vLLM

PVC is used to store the model cache and it is optional, you can use hostPath or other storage options

```
```yaml
```

```
apiVersion: v1
```

```
kind: PersistentVolumeClaim
```

metadata:

name: mistral-7b

namespace: default

spec:

accessModes:

- ReadWriteOnce

resources:

requests:

storage: 50Gi

storageClassName: default

volumeMode: Filesystem

...

Secret is optional and only required for accessing gated models, you can skip this step if you are not using gated models

```
```yaml
```

apiVersion: v1

kind: Secret

metadata:

name: hf-token-secret

namespace: default

type: Opaque

stringData:

token: "REPLACE\_WITH\_TOKEN"

...

Next to create the deployment file for vLLM to run the model server. The following example deploys the `Mistral-7B-Instruct-v0.3` model.

Here are two examples for using NVIDIA GPU and AMD GPU.

NVIDIA GPU:

```
```yaml
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: mistral-7b
```

```
  namespace: default
```

```
  labels:
```

```
    app: mistral-7b
```

```
spec:
```

```
  replicas: 1
```

```
  selector:
```

```
    matchLabels:
```

```
      app: mistral-7b
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: mistral-7b
```

```
    spec:
```

```
      volumes:
```

```
        - name: cache-volume
```

persistentVolumeClaim:

claimName: mistral-7b

# vLLM needs to access the host's shared memory for tensor parallel inference.

- name: shm

emptyDir:

medium: Memory

sizeLimit: "2Gi"

containers:

- name: mistral-7b

image: vllm/vllm-openai:latest

command: ["/bin/sh", "-c"]

args: [

"vllm serve mistralai/Mistral-7B-Instruct-v0.3 --trust-remote-code --enable-chunked-prefill

--max\_num\_batched\_tokens 1024"

]

env:

- name: HUGGING\_FACE\_HUB\_TOKEN

valueFrom:

secretKeyRef:

name: hf-token-secret

key: token

ports:

- containerPort: 8000

resources:

limits:

cpu: "10"

memory: 20G

nvidia.com/gpu: "1"

requests:

cpu: "2"

memory: 6G

nvidia.com/gpu: "1"

volumeMounts:

- mountPath: /root/.cache/huggingface

name: cache-volume

- name: shm

mountPath: /dev/shm

livenessProbe:

httpGet:

path: /health

port: 8000

initialDelaySeconds: 60

periodSeconds: 10

readinessProbe:

httpGet:

path: /health

port: 8000

initialDelaySeconds: 60

periodSeconds: 5

...

AMD GPU:

You can refer to the `deployment.yaml` below if using AMD ROCm GPU like MI300X.

```
```yaml
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: mistral-7b
```

```
  namespace: default
```

```
  labels:
```

```
    app: mistral-7b
```

```
spec:
```

```
  replicas: 1
```

```
  selector:
```

```
    matchLabels:
```

```
      app: mistral-7b
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: mistral-7b
```

```
    spec:
```

```
      volumes:
```

```
        # PVC
```

```
        - name: cache-volume
```

```
          persistentVolumeClaim:
```

```
            claimName: mistral-7b
```

```
        # vLLM needs to access the host's shared memory for tensor parallel inference.
```

```
        - name: shm
```

```
          emptyDir:
```

medium: Memory

sizeLimit: "8Gi"

hostNetwork: true

hostIPC: true

containers:

- name: mistral-7b

image: rocm/vllm:rocm6.2\_mi300\_ubuntu20.04\_py3.9\_vllm\_0.6.4

securityContext:

seccompProfile:

type: Unconfined

runAsGroup: 44

capabilities:

add:

- SYS\_PTRACE

command: ["/bin/sh", "-c"]

args: [

"vllm serve mistralai/Mistral-7B-v0.3 --port 8000 --trust-remote-code

--enable-chunked-prefill --max\_num\_batched\_tokens 1024"

]

env:

- name: HUGGING\_FACE\_HUB\_TOKEN

valueFrom:

secretKeyRef:

name: hf-token-secret

key: token

ports:

- containerPort: 8000

resources:

limits:

cpu: "10"

memory: 20G

amd.com/gpu: "1"

requests:

cpu: "6"

memory: 6G

amd.com/gpu: "1"

volumeMounts:

- name: cache-volume

mountPath: /root/.cache/huggingface

- name: shm

mountPath: /dev/shm

...

You can get the full example with steps and sample yaml files from <https://github.com/ROCm/k8s-device-plugin/tree/master/example/vllm-serve>.

## 2. Create a Kubernetes Service for vLLM

Next, create a Kubernetes Service file to expose the `mistral-7b` deployment:

```
```yaml
```

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```



```
name: mistral-7b

namespace: default

spec:

  ports:

    - name: http-mistral-7b

      port: 80

      protocol: TCP

      targetPort: 8000

    # The label selector should match the deployment labels & it is useful for prefix caching feature

  selector:

    app: mistral-7b

  sessionAffinity: None

  type: ClusterIP

...

```

### 3. Deploy and Test

Apply the deployment and service configurations using `kubectl apply -f <filename>`:

```
```console

kubectl apply -f deployment.yaml

kubectl apply -f service.yaml

...

```

To test the deployment, run the following `curl` command:

```
```console

```

```
curl http://mistral-7b.default.svc.cluster.local/v1/completions \
-H "Content-Type: application/json" \
-d '{
    "model": "mistralai/Mistral-7B-Instruct-v0.3",
    "prompt": "San Francisco is a",
    "max_tokens": 7,
    "temperature": 0
}'
...
```

If the service is correctly deployed, you should receive a response from the vLLM model.

## ## Conclusion

Deploying vLLM with Kubernetes allows for efficient scaling and management of ML models leveraging GPU resources. By following the steps outlined above, you should be able to set up and test a vLLM deployment within your Kubernetes cluster. If you encounter any issues or have suggestions, please feel free to contribute to the documentation.