

(nginxloadbalancer)=

## # Using Nginx

This document shows how to launch multiple vLLM serving containers and use Nginx to act as a load balancer between the servers.

Table of contents:

1. [Build Nginx Container](#nginxloadbalancer-nginx-build)
2. [Create Simple Nginx Config file](#nginxloadbalancer-nginx-conf)
3. [Build vLLM Container](#nginxloadbalancer-nginx-vllm-container)
4. [Create Docker Network](#nginxloadbalancer-nginx-docker-network)
5. [Launch vLLM Containers](#nginxloadbalancer-nginx-launch-container)
6. [Launch Nginx](#nginxloadbalancer-nginx-launch-nginx)
7. [Verify That vLLM Servers Are Ready](#nginxloadbalancer-nginx-verify-nginx)

(nginxloadbalancer-nginx-build)=

## ## Build Nginx Container

This guide assumes that you have just cloned the vLLM project and you're currently in the vllm root directory.

```
```console
```

```
export vllm_root=`pwd`
```

```
```
```

Create a file named `Dockerfile.nginx`:

```
```console
```

```
FROM nginx:latest
```

```
RUN rm /etc/nginx/conf.d/default.conf
```

```
EXPOSE 80
```

```
CMD ["nginx", "-g", "daemon off;"]
```

```
```
```

Build the container:

```
```console
```

```
docker build . -f Dockerfile.nginx --tag nginx-lb
```

```
```
```

(nginxloadbalancer-nginx-conf)=

## Create Simple Nginx Config file

Create a file named `nginx\_conf/nginx.conf`. Note that you can add as many servers as you'd like.

In the below example we'll start with two. To add more, add another `server vllmN:8000 max\_fails=3

fail\_timeout=10000s;` entry to `upstream backend`.

```
```console
```

```
upstream backend {
```

```
    least_conn;
```

```

server vllm0:8000 max_fails=3 fail_timeout=10000s;

server vllm1:8000 max_fails=3 fail_timeout=10000s;

}

server {

    listen 80;

    location / {

        proxy_pass http://backend;

        proxy_set_header Host $host;

        proxy_set_header X-Real-IP $remote_addr;

        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;

        proxy_set_header X-Forwarded-Proto $scheme;

    }

}

...

```

(nginxloadbalancer-nginx-vllm-container)=

## Build vLLM Container

```

```console

cd $vllm_root

docker build -f Dockerfile . --tag vllm

...

```

If you are behind proxy, you can pass the proxy settings to the docker build command as shown below:

```
```console
```

```
cd $vllm_root
```

```
docker build -f Dockerfile . --tag vllm --build-arg http_proxy=$http_proxy --build-arg
```

```
https_proxy=$https_proxy
```

```
```
```

```
(nginxloadbalancer-nginx-docker-network)=
```

```
## Create Docker Network
```

```
```console
```

```
docker network create vllm_nginx
```

```
```
```

```
(nginxloadbalancer-nginx-launch-container)=
```

```
## Launch vLLM Containers
```

Notes:

- If you have your HuggingFace models cached somewhere else, update `hf\_cache\_dir` below.
- If you don't have an existing HuggingFace cache you will want to start `vllm0` and wait for the model to complete downloading and the server to be ready. This will ensure that `vllm1` can leverage the model you just downloaded and it won't have to be downloaded again.
- The below example assumes GPU backend used. If you are using CPU backend, remove `--gpus all`, add `VLLM\_CPU\_KVCACHE\_SPACE` and `VLLM\_CPU\_OMP\_THREADS\_BIND` environment variables to the docker run command.

- Adjust the model name that you want to use in your vLLM servers if you don't want to use `Llama-2-7b-chat-hf`.

```console

```
mkdir -p ~/.cache/huggingface/hub/
```

```
hf_cache_dir=~/.cache/huggingface/
```

```
docker run -itd --ipc host --privileged --network vllm_nginx --gpus all --shm-size=10.24gb -v  
$hf_cache_dir:/root/.cache/huggingface/ -p 8081:8000 --name vllm0 vllm --model  
meta-llama/Llama-2-7b-chat-hf
```

```
docker run -itd --ipc host --privileged --network vllm_nginx --gpus all --shm-size=10.24gb -v  
$hf_cache_dir:/root/.cache/huggingface/ -p 8082:8000 --name vllm1 vllm --model  
meta-llama/Llama-2-7b-chat-hf
```

```

:::{note}

If you are behind proxy, you can pass the proxy settings to the docker run command via `-e http\_proxy=\$http\_proxy -e https\_proxy=\$https\_proxy`.

:::

(nginxloadbalancer-nginx-launch-nginx)=

## Launch Nginx

```console

```
docker run -itd -p 8000:80 --network vllm_nginx -v ./nginx_conf:/etc/nginx/conf.d/ --name nginx-lb  
nginx-lb:latest
```

```

(nginxloadbalancer-nginx-verify-nginx)=

## Verify That vLLM Servers Are Ready

```console

docker logs vllm0 | grep Uvicorn

docker logs vllm1 | grep Uvicorn

```

Both outputs should look like this:

```console

INFO: Uvicorn running on http://0.0.0.0:8000 (Press CTRL+C to quit)

```